



# Advanced Network Security

Amir Mahdi Sadeghzadeh, Ph.D.

# Outline

- What is network security?
- Principles of cryptography
- **Authentication, message integrity**
- Securing e-mail
- Securing TCP connections: TLS
- Network layer security: IPsec



# Other than Confidentiality

## ■ Authentication

- Definition: The process of verifying the identity of a user, device, or system.
  - Ensures that the entity accessing the system is who they claim to be.
- Methods: Passwords, biometrics, digital signatures, MACs, etc.

## ■ Data Integrity

- Definition: The assurance that data remains unaltered during transmission or storage.
- Methods: MACs, digital signatures.

## ■ Non-repudiation

- Definition: The inability of a user to deny the authenticity or origin of a communication or action.
- Methods: Digital signatures.

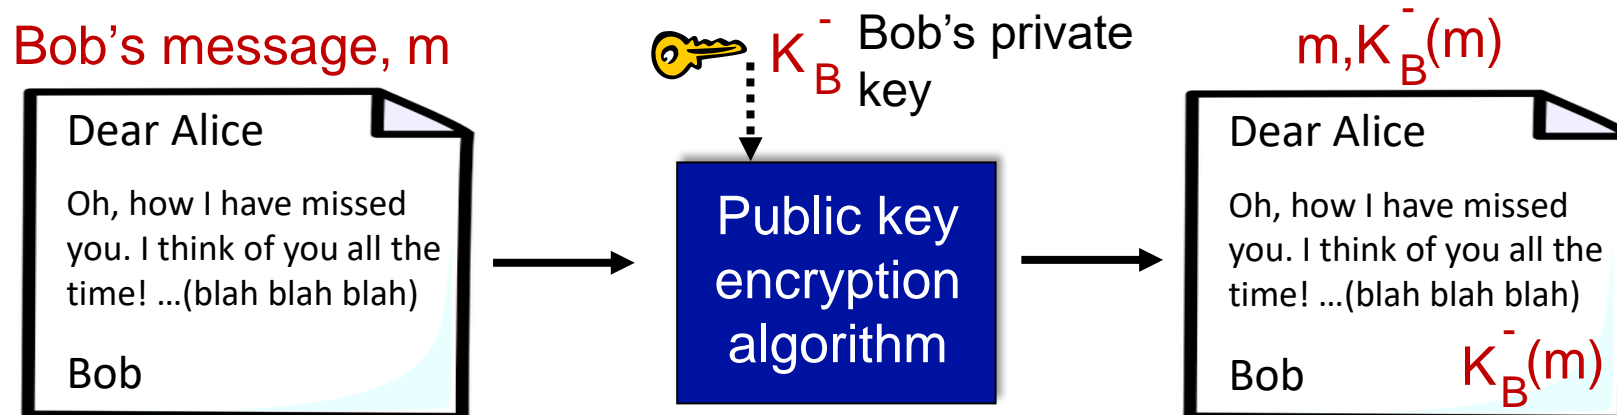
# Where Does This Fit?

	Secret Key Setting	Public Key Setting
Secrecy / Confidentiality	Stream cipher Block cipher + encryption modes	Public key encryption: RSA, El Gamal, etc.
Authenticity / Integrity	MAC	Digital Signatures

# Digital signatures


## Cryptographic technique analogous to hand-written signatures:

- sender (Bob) digitally signs document: he is document owner/creator.
- *verifiable, nonforgeable*: recipient (Alice) can prove to someone that Bob, and no one else (including Alice), must have signed document
- **simple digital signature for message  $m$ :**
  - Bob signs  $m$  by encrypting with his private key  $K_B$ , creating “signed” message,  $K_B^-(m)$



# Message digests

- Computationally expensive to public-key-encrypt long messages
  - fixed-length, easy- to-compute digital “fingerprint”
  - apply hash function  $H$  to  $m$ , get fixed size message digest,  $H(m)$
- Existential forgery
  - extra layer of security

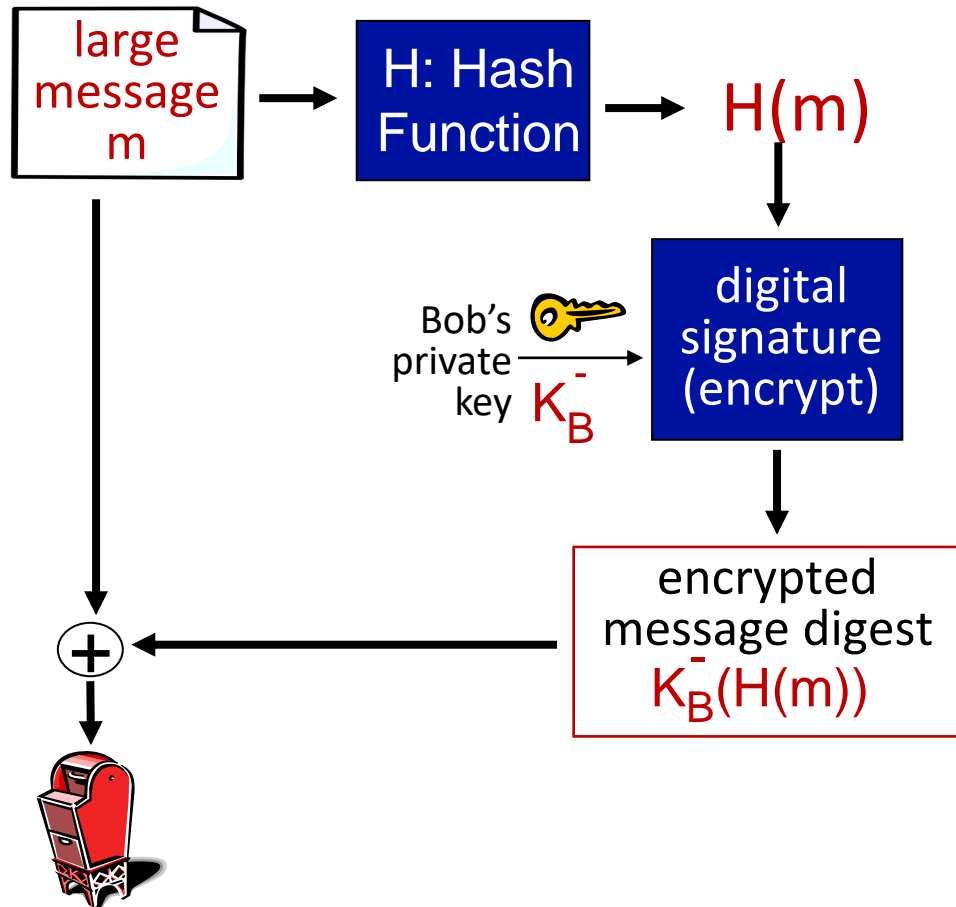


## Hash function properties:

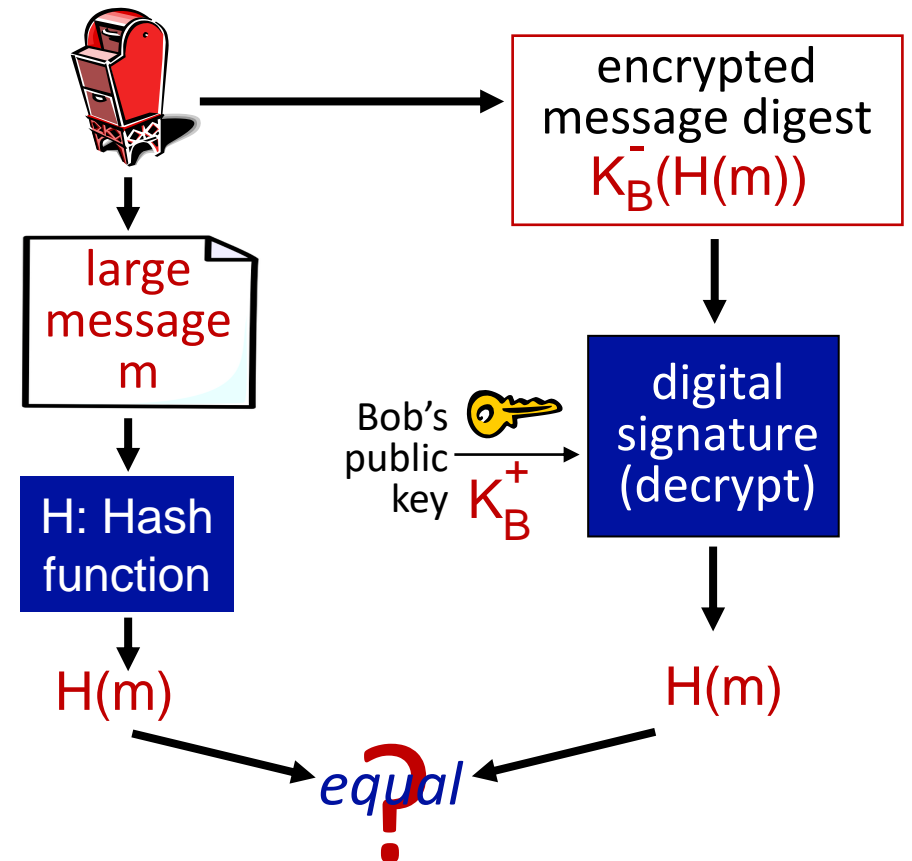
- produces fixed-size msg digest (fingerprint)
- given message digest  $x$ , computationally infeasible to find  $m$  such that  $x = H(m)$

# Digital signature = signed message digest


Bob sends digitally signed message:



Alice verifies signature, integrity of digitally signed message:



# Message Authentication Codes (MACs)



**Generate tag:**

$$\text{tag} \leftarrow S(k, m)$$

**Verify tag:**

$$V(k, m, \text{tag}) \stackrel{?}{=} \text{'yes'}$$

Def: **MAC**  $I = (S, V)$  defined over  $(K, M, T)$  is a pair of algs:

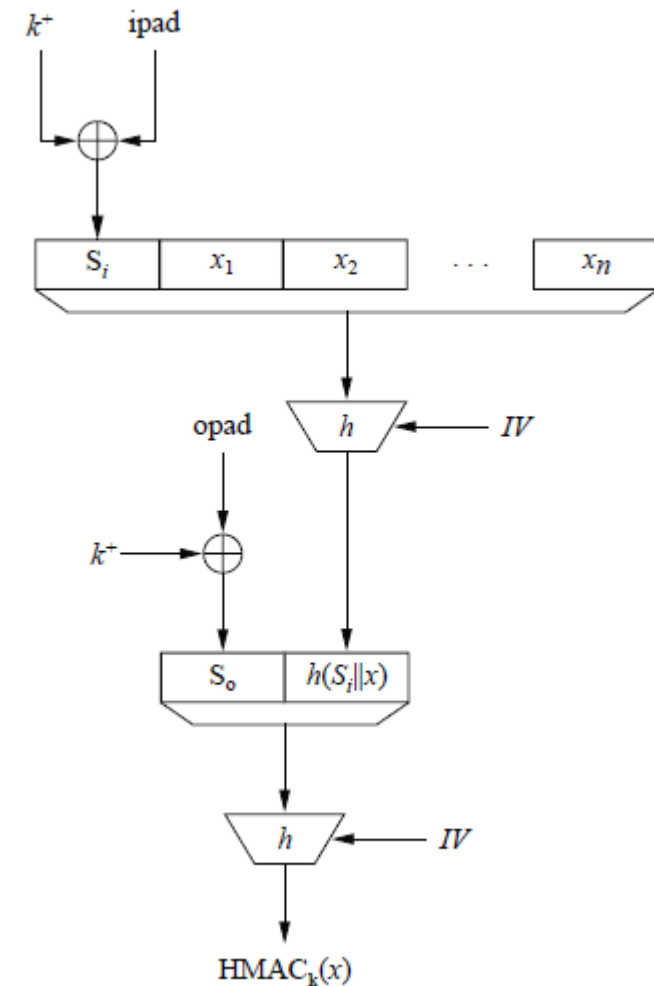
- $S(k, m)$  outputs  $t$  in  $T$
- $V(k, m, t)$  outputs 'yes' or 'no'



# HMAC

- A hash-based message authentication code which does not show the security weakness of prefix and suffix MACs construction
  - The scheme consists of an inner and outer hash.

$$HMAC_k(x) = h[(k^+ \oplus opad) || h[(k^+ \oplus ipad) || x]]$$



# Authentication Protocol

**Goal:** Bob wants Alice to “prove” her identity to him

**Protocol ap1.0:** Alice says “I am Alice”




*failure scenario??*



# Authentication Protocol

**Goal:** Bob wants Alice to “prove” her identity to him

**Protocol ap1.0:** Alice says “I am Alice”




*in a network, Bob  
can not “see”  
Alice, so Trudy  
simply declares  
herself to be Alice*



# Authentication: another try

**Goal:** Bob wants Alice to “prove” her identity to him

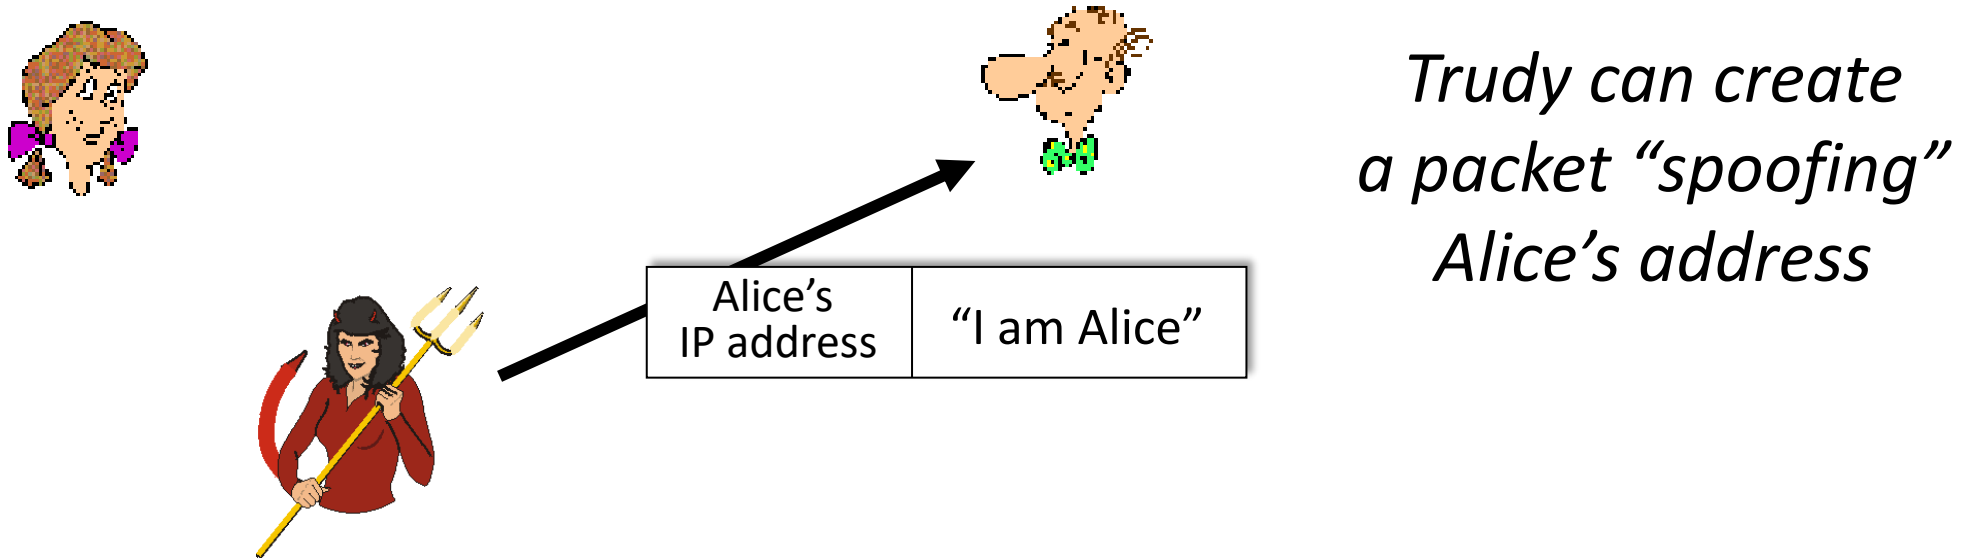
**Protocol ap2.0:** Alice says “I am Alice” in an IP packet containing her source IP address



# Authentication: another try

**Goal:** Bob wants Alice to “prove” her identity to him

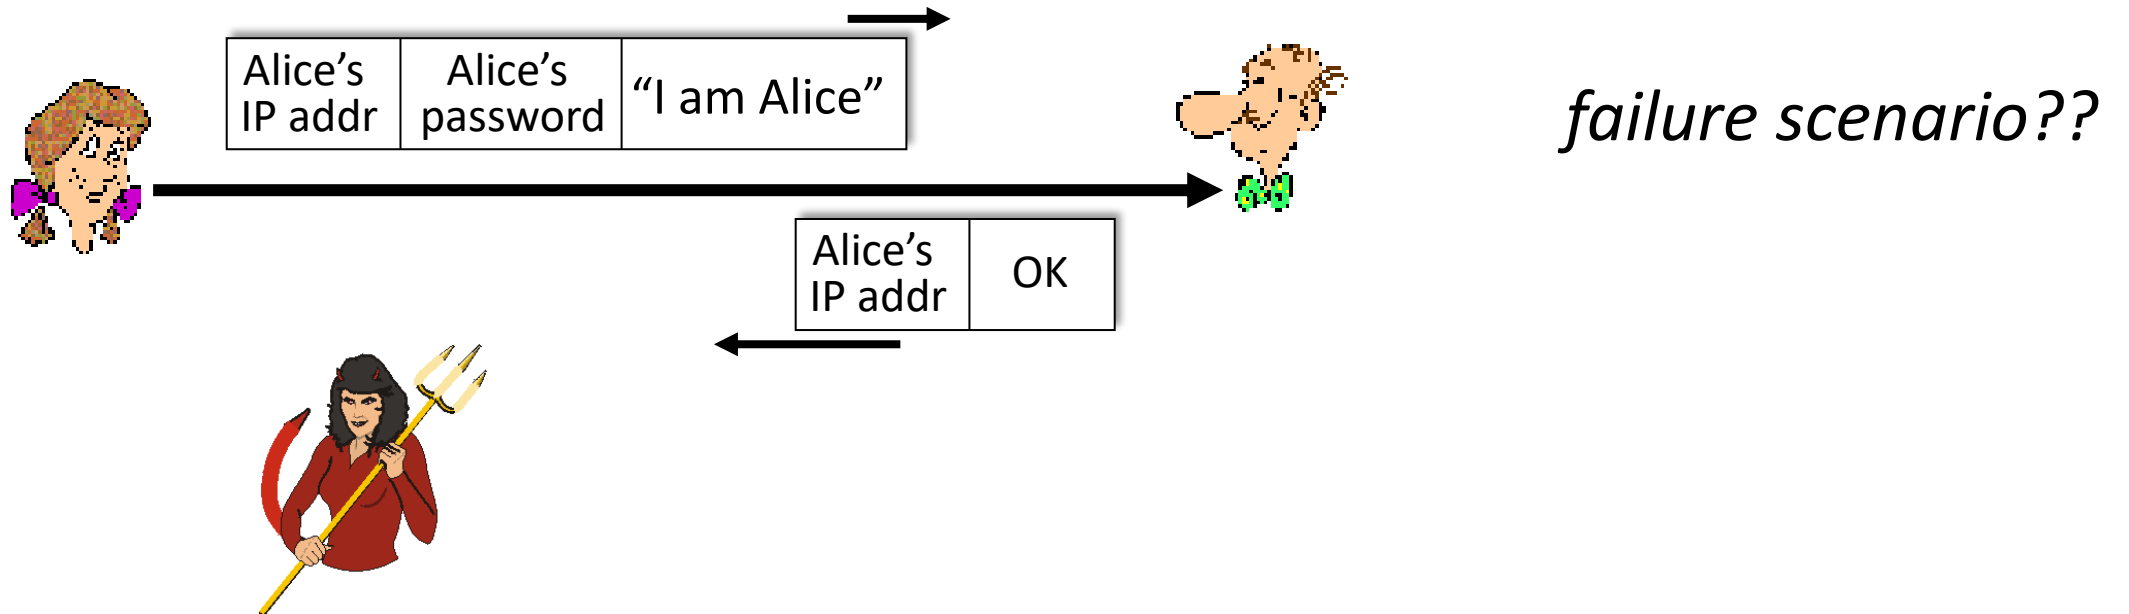
**Protocol ap2.0:** Alice says “I am Alice” in an IP packet containing her source IP address



# Authentication: a third try

**Goal:** Bob wants Alice to “prove” her identity to him

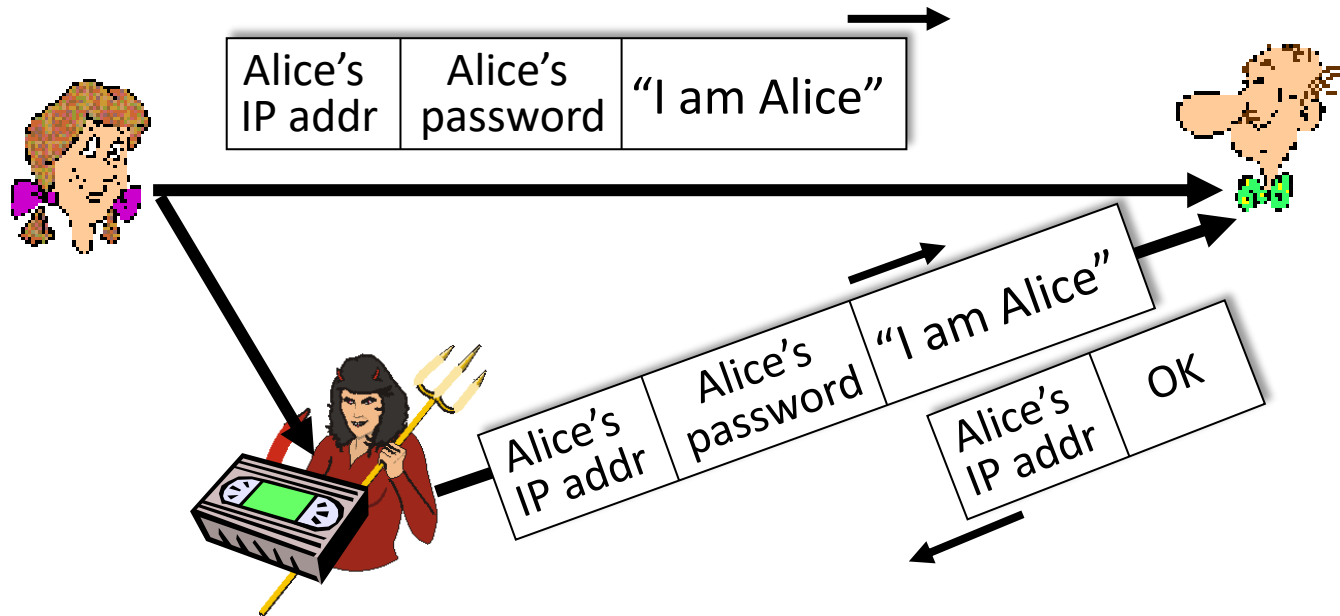
**Protocol ap3.0:** Alice says “I am Alice” and sends her secret password to “prove” it.



# Authentication: a third try

**Goal:** Bob wants Alice to “prove” her identity to him

**Protocol ap3.0:** Alice says “I am Alice” and sends her secret password to “prove” it.

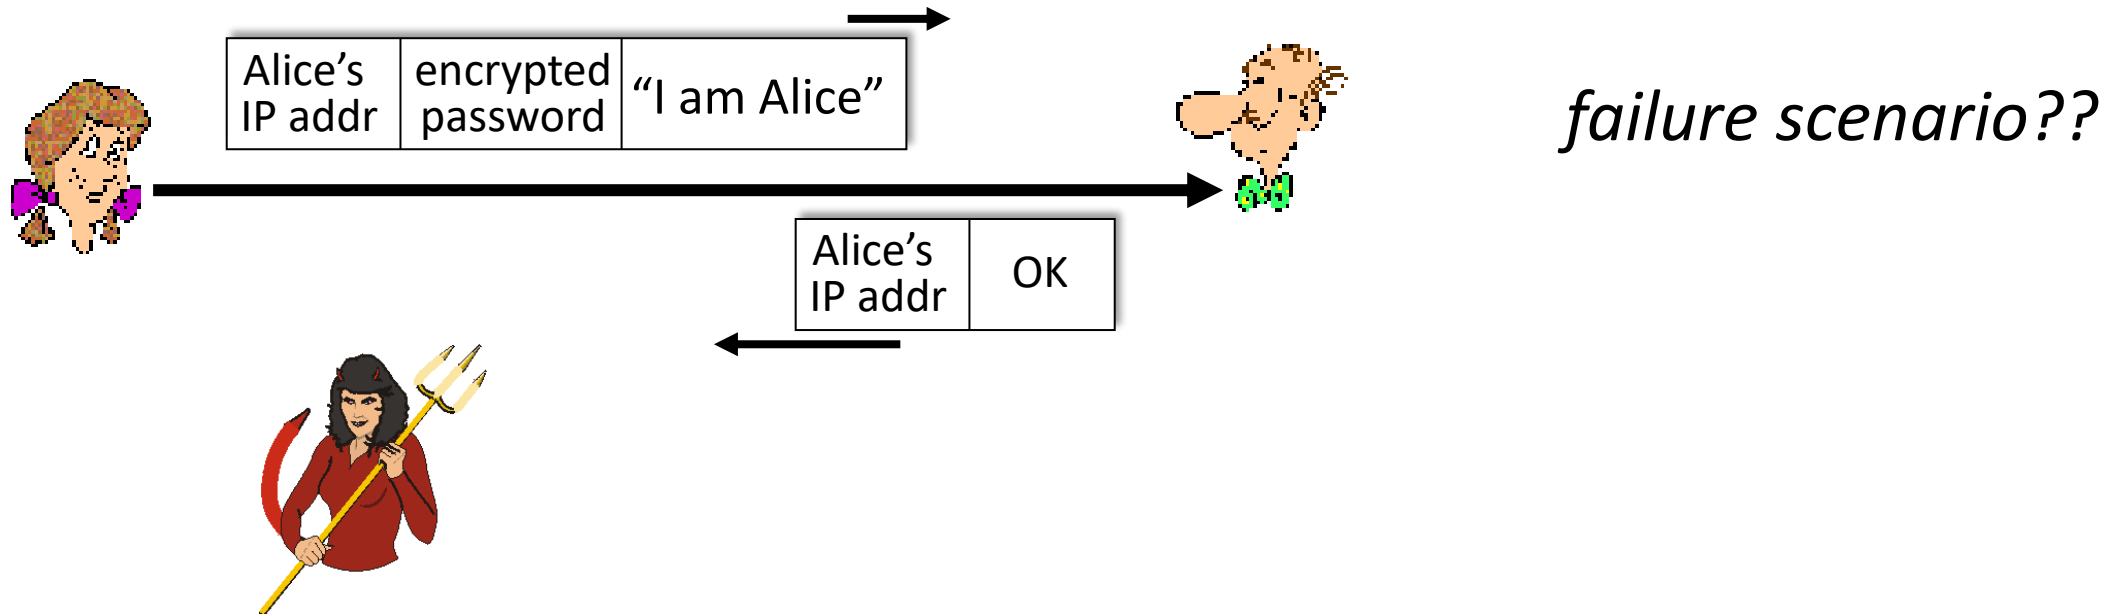


*playback attack:  
Trudy records  
Alice's packet  
and later  
plays it back to Bob*

# Authentication: a modified third try

**Goal:** Bob wants Alice to “prove” her identity to him

**Protocol ap3.0:** Alice says “I am Alice” and sends her encrypted secret password to “prove” it.

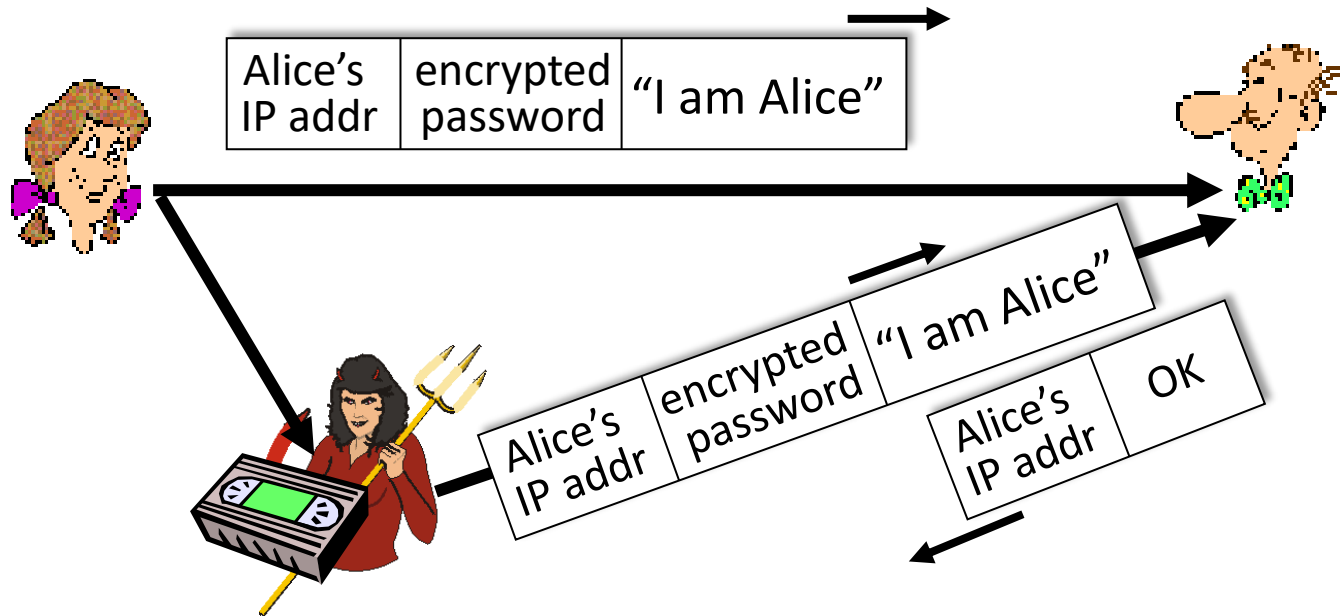




# Authentication: a modified third try

**Goal:** Bob wants Alice to “prove” her identity to him

**Protocol ap3.0:** Alice says “I am Alice” and sends her encrypted secret password to “prove” it.



*playback attack still works: Trudy records Alice's packet and later plays it back to Bob*

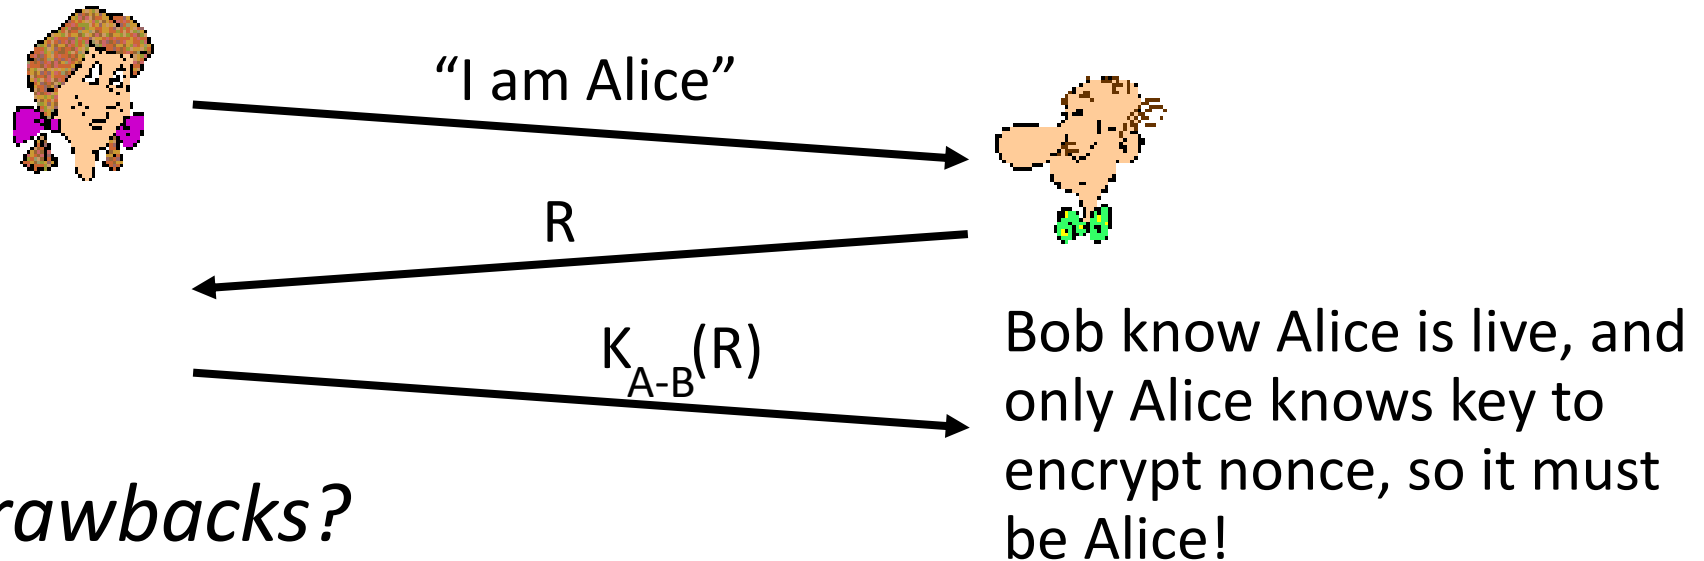
# Authentication: a fourth try

**Goal:** avoid playback attack

**nonce:** number (R) used only **once-in-a-lifetime**

**protocol ap4.0:** to prove Alice “live”, Bob sends Alice nonce, R

- Alice must return R, encrypted with shared secret key

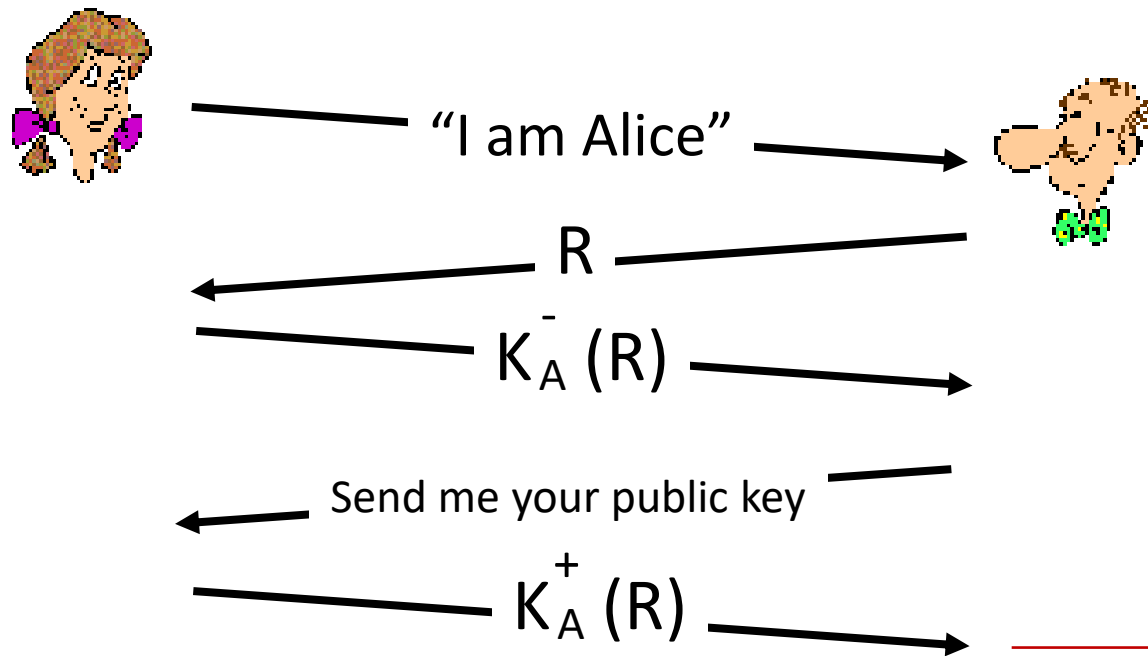


*Failures, drawbacks?*

# Authentication: ap5.0

ap4.0 requires shared symmetric key - can we authenticate using public key techniques?

**ap5.0:** use nonce, public key cryptography



Bob computes

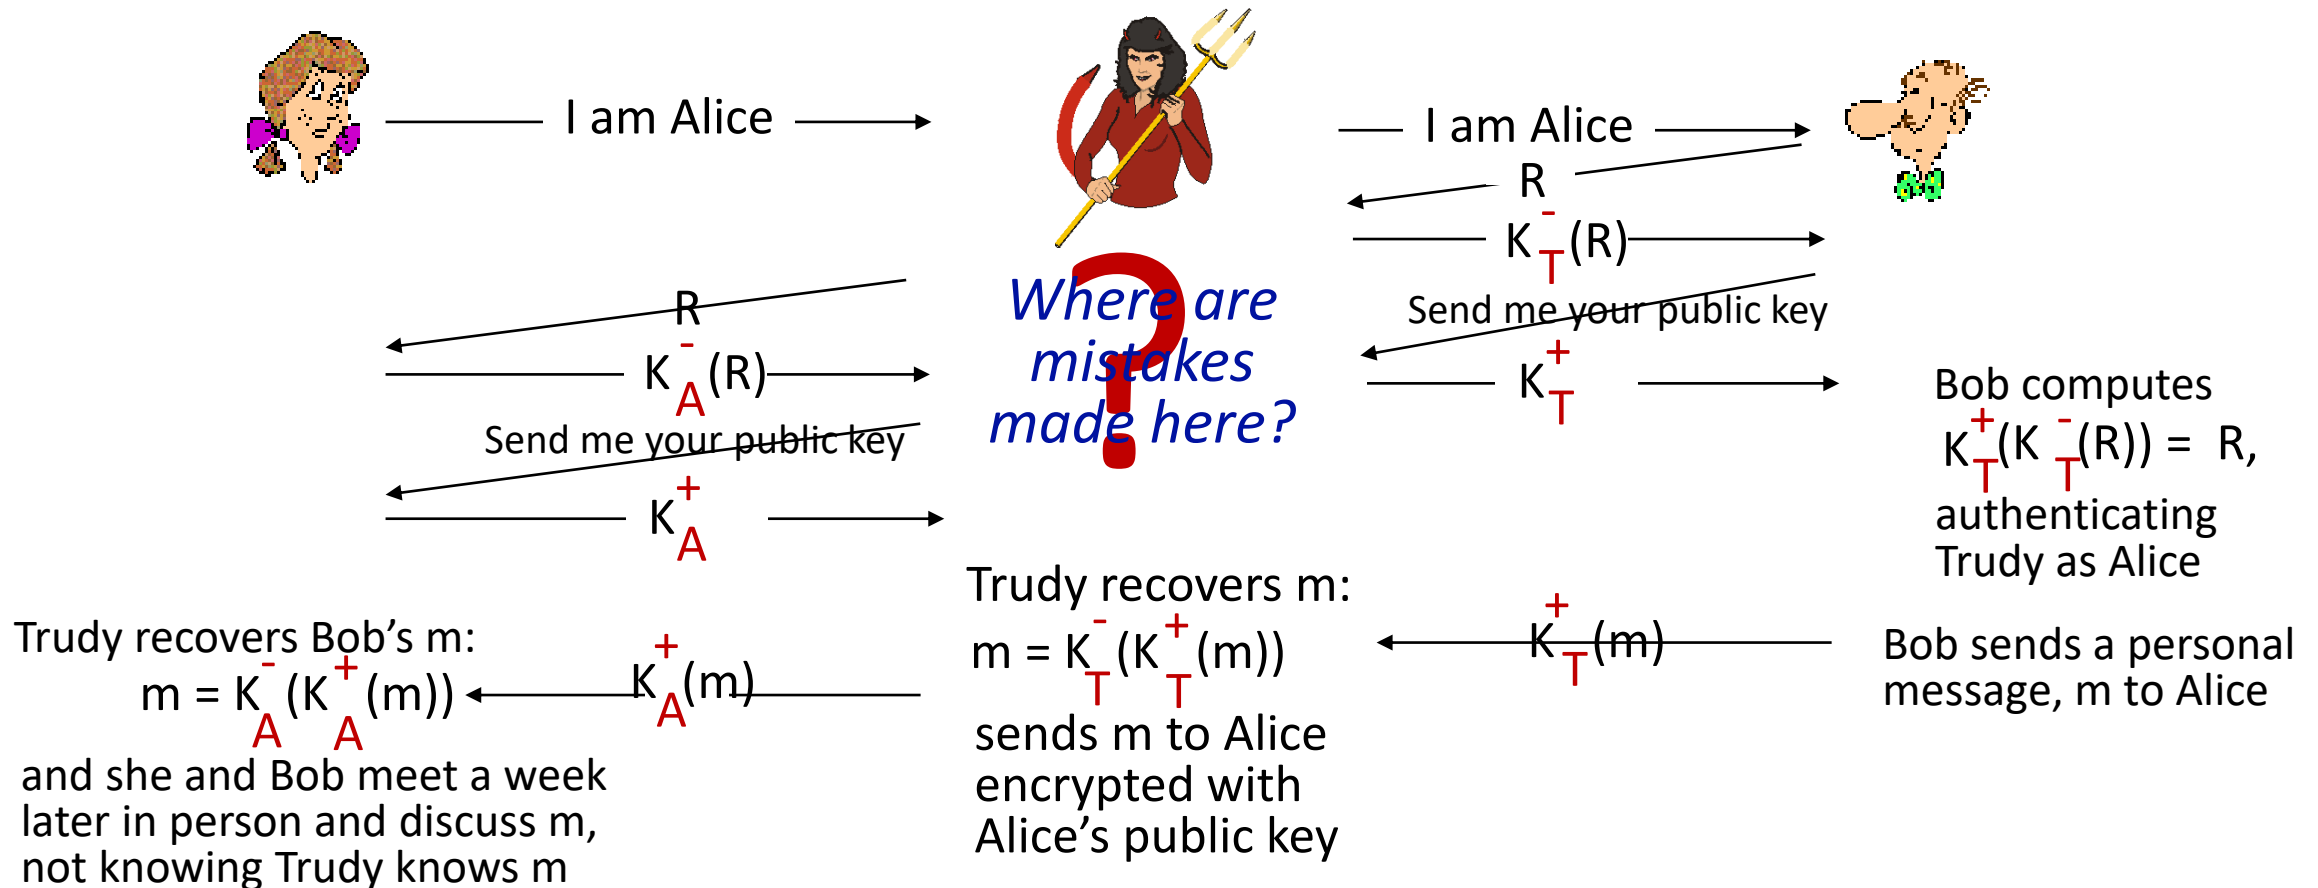
$$K_A^+ (K_A^-(R)) = R$$

and knows only Alice could have the private key, that encrypted  $R$  such that

$$K_A^+ (K_A^-(R)) = R$$

# Authentication: ap5.0 – there's still a flaw!

**man (or woman) in the middle attack:** Trudy poses as Alice (to Bob) and as Bob (to Alice)



# Need for certified public keys

- motivation: Trudy plays pizza prank on Bob
  - Trudy creates e-mail order:  
*Dear Pizza Store, Please deliver to me four pepperoni pizzas. Thank you, Bob*
  - Trudy signs order with her private key
  - Trudy sends order to Pizza Store
  - Trudy sends to Pizza Store her public key, but says it's Bob's public key
  - Pizza Store verifies signature; then delivers four pepperoni pizzas to Bob
  - Bob doesn't even like pepperoni



# Public key Certification Authorities (CA)

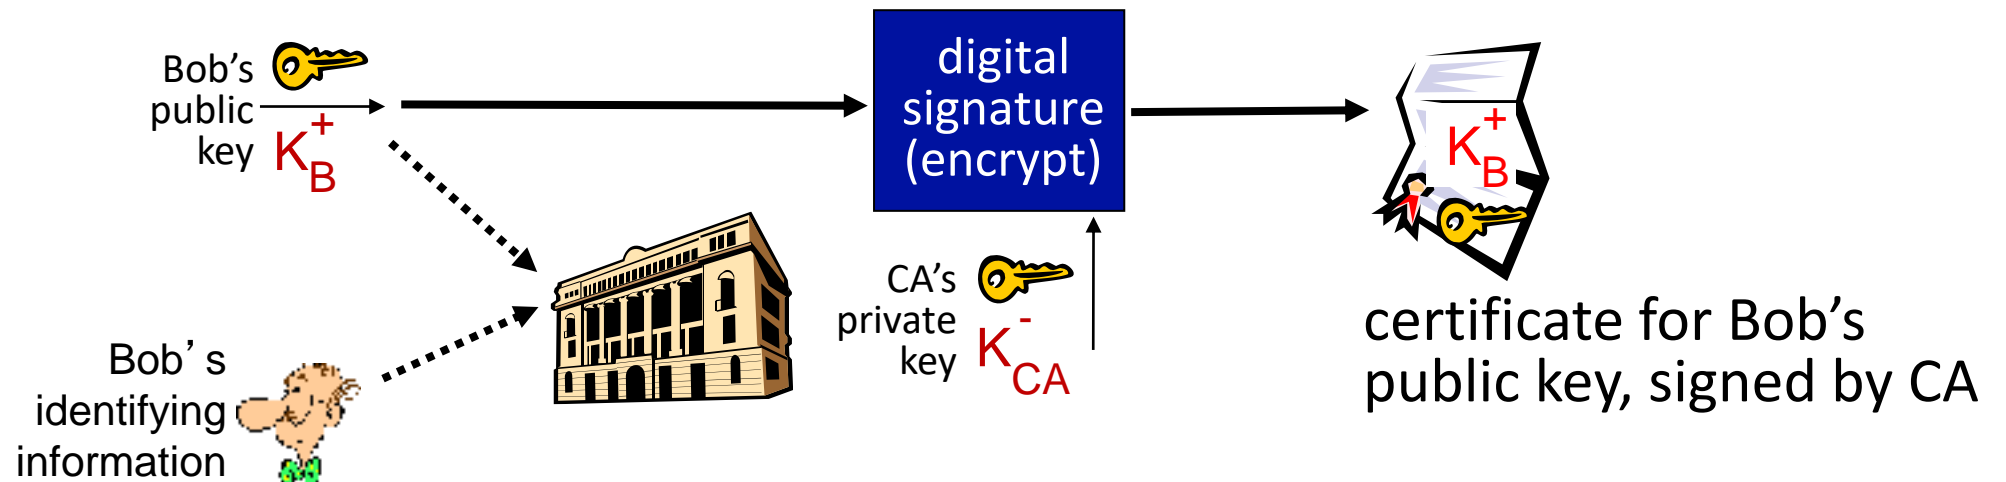
- **certification authority (CA):** binds public key to particular entity, E
- entity (person, website, router) registers its public key with “proof of identity” to CA
  - CA creates certificate binding identity E to E’s public key

# Certification Scheme

- We can place the following requirements on Certificate scheme
  - **Any participant can read a certificate** to determine the **name** and **public key** of the certificate's owner.
  - **Any participant can verify** that the certificate originated from the certificate authority and is not counterfeit.
  - **Only the certificate authority** can **create** and **update** certificates.
  - Any participant can **verify the time validity** of the certificate.

# Public Key Certification

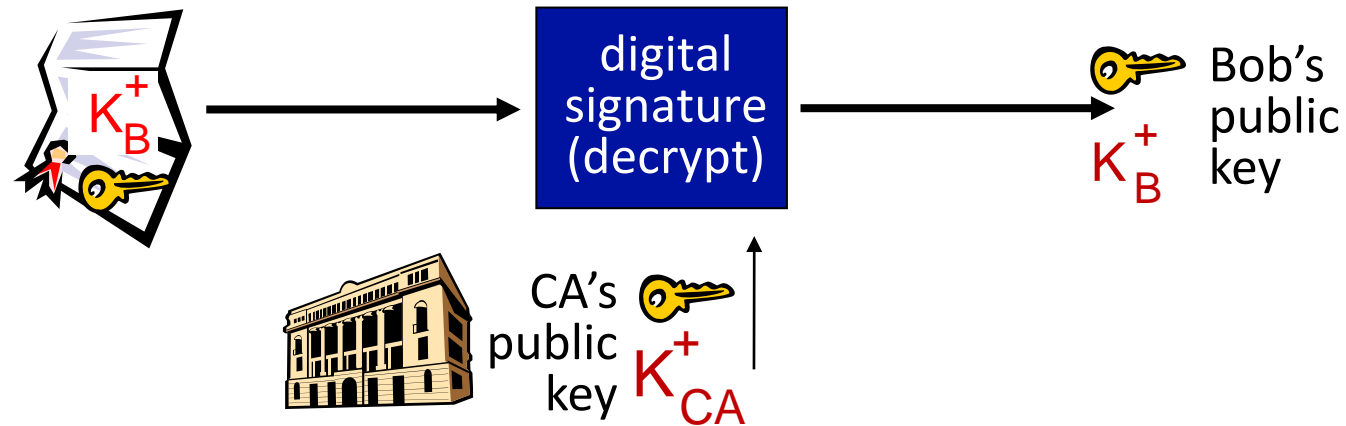
- Certificate containing E's public key digitally signed by CA: CA says "this is E's public key"





# Public key Certification Authorities (CA)

- when Alice wants Bob's public key:
  - gets Bob's certificate (Bob or elsewhere)
  - apply CA's public key to Bob's certificate, get Bob's public key



# Certification Scheme

- For participant A, the **authority provides a certificate** of the form

$$C_A = E(PR_{auth}, [T || ID_A || PU_A])$$

where  $PR_{auth}$  is the private key used by the authority and T is a timestamp.  $PU_a$  and  $ID_A$  are the Alice's public key and ID.

- **Anyone can verify** the certificate as follows

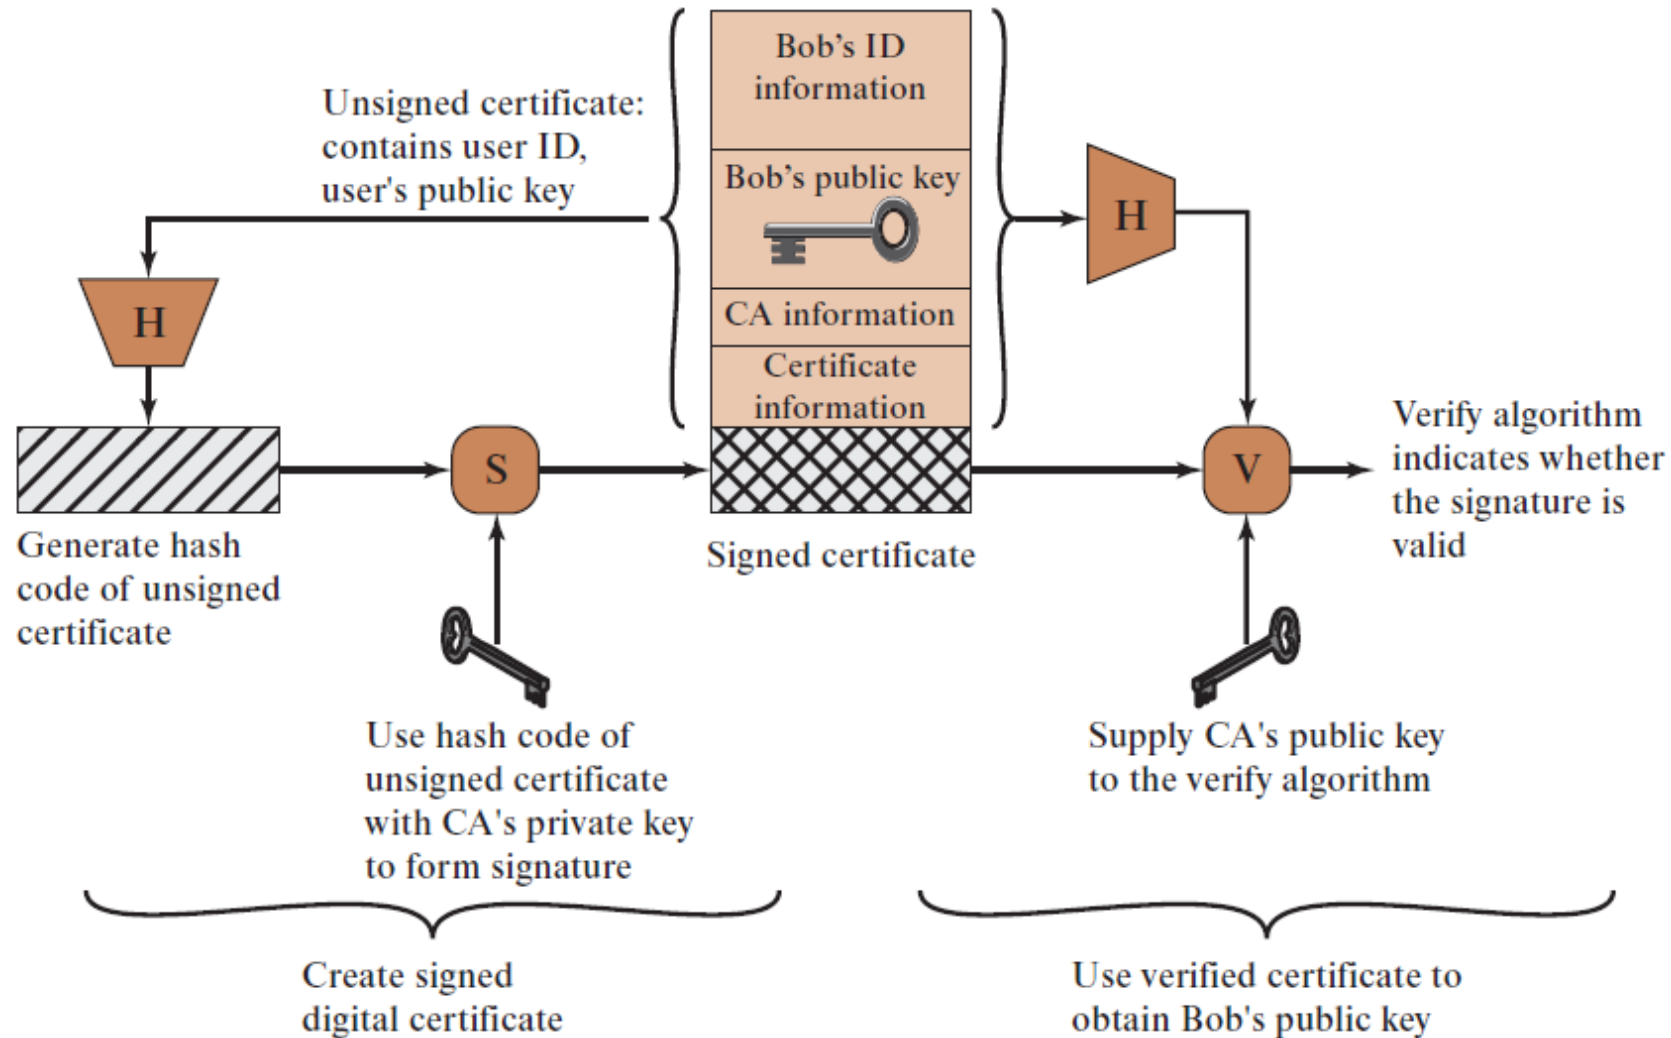
$$D(PU_{auth}, C_A) = D(PU_{auth}, E(PR_{auth}, [T || ID_A || PU_A])) = (T || ID_A || PU_A)$$

where  $PU_{auth}$  is the public key used by the authority

# X.509 Certificates

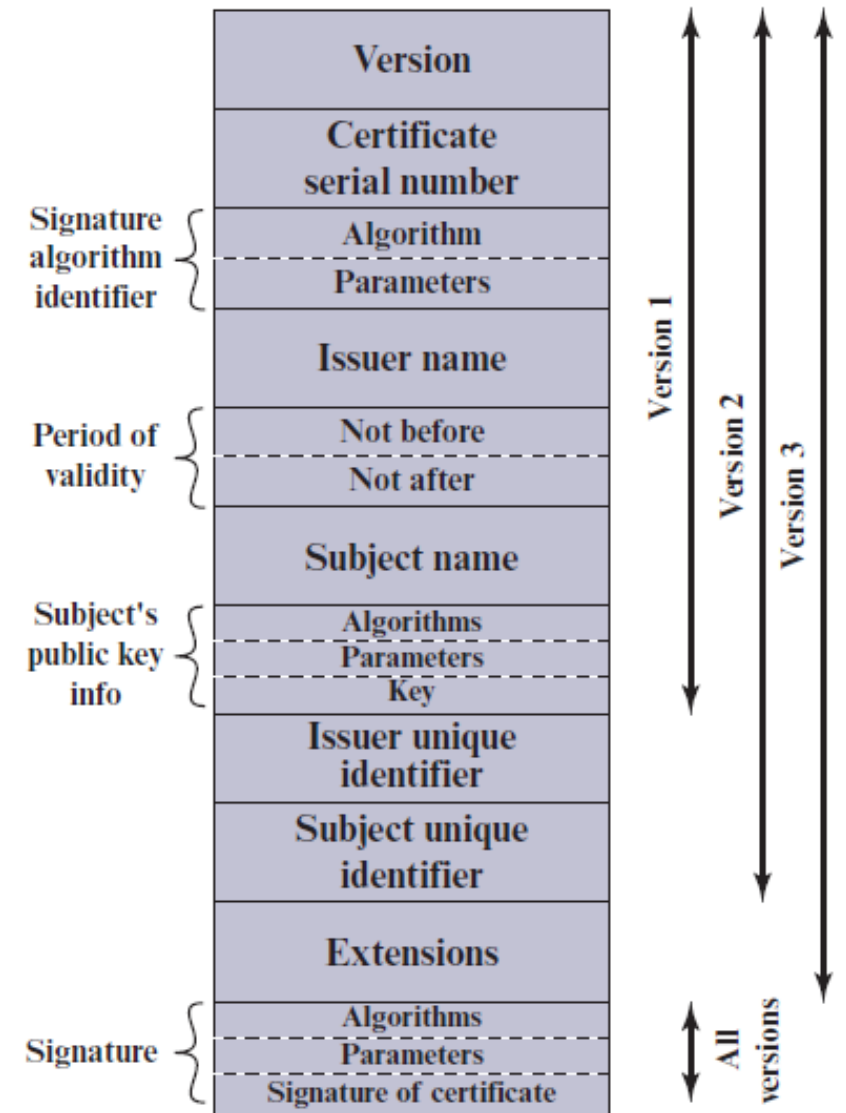
- X.509 is **the Internationally accepted standard** for how to construct a public key certificate, and is becoming widely used.
  - It has gone through several versions.
- Defines framework for authentication services:
  - Defines that public keys stored as certificates in a public directory.
  - Certificates are issued and signed by an entity called certification authority (CA).
  - Used by numerous applications: SSL, IPSec, SET
  - Example: see certificates accepted by your browser

# X.509 Certificates



# X.509 Certificates

- issued by a Certification Authority (CA), containing:
  - version V (1, 2, or 3)
  - serial number SN (unique within CA) identifying certificate
  - signature algorithm identifier
  - issuer name
  - period of validity (from - to dates)
  - subject name (name of owner)
  - subject public-key info Ap (algorithm, parameters, key)
  - issuer unique identifier
  - subject unique identifier
  - extension fields
  - signature (of hash of all fields in certificate)



(a) X.509 certificate

Figure 15.11 X.509 Formats

# Verifying Identity

- A CA verifies some site's identity. How?
- It's easy to verify the existence of a corporation—but how do you verify who speaks for it?
- For publicly traded companies, top executives are generally a matter of public record—do the CEO or CTO have to show up in person?

# Many Different Ways!

- From the Baseline Requirements document from the CA Browser Forum, verification can be done by:
  - Documents from or communication with a government agency;
  - An in-person visit by the CA;
  - A reliable third-party database;
  - An “attestation letter” from a lawyer or accountant, etc.;
  - More. . .
- These procedures are manual, annoying, and complex—can we automate them?

# Automated Web CAs

- Web site certificates should be issued to the owner of the web site, e.g., the party that controls the site
- Demonstrate control of the web site by doing something only the site owner can do
- Example: the CA sends the site a random number; the site puts that number into a specific URL




# Outline

- What is network security?
- Principles of cryptography
- Authentication, message integrity
- **Securing e-mail**
- Securing TCP connections: TLS
- Network layer security: IPsec



# Secure e-mail: confidentiality

Alice wants to send *confidential* e-mail,  $m$ , to Bob.




Alice:

- generates random *symmetric* private key,  $K_S$
- encrypts message with  $K_S$  (for efficiency)
- also encrypts  $K_S$  with Bob's public key
- sends both  $K_S(m)$  and  $K_B^+(K_S)$  to Bob

# Secure e-mail: confidentiality (more)

Alice wants to send *confidential* e-mail,  $m$ , to Bob.




**Bob:**

- uses his private key to decrypt and recover  $K_S$
- uses  $K_S$  to decrypt  $K_S(m)$  to recover  $m$

# Secure e-mail: integrity, authentication

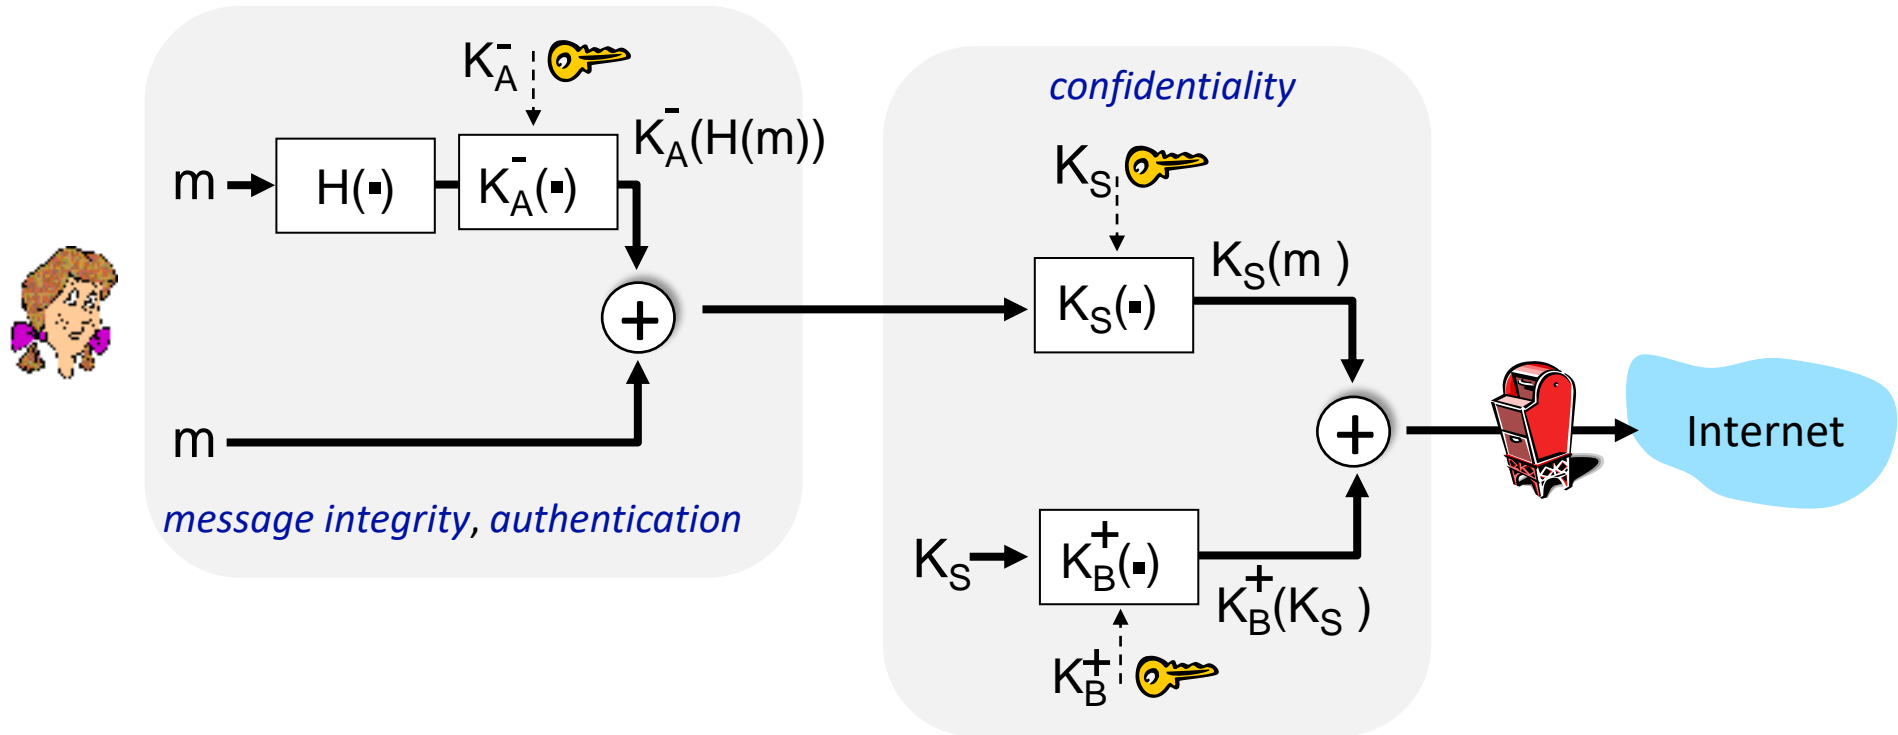
Alice wants to send  $m$  to Bob, with *message integrity, authentication*



- Alice digitally signs hash of her message with her private key, providing integrity and authentication
- sends both message (in the clear) and digital signature

# Secure e-mail: confidentiality, integrity, authen.

Alice sends  $m$  to Bob, with *confidentiality, message integrity, authentication*



**Alice uses three keys:** her private key, Bob's public key, new symmetric key

*What are Bob's complementary actions?*

# Outline

- What is network security?
- Principles of cryptography
- Authentication, message integrity
- Securing e-mail
- **Securing TCP connections: TLS**
- Network layer security: IPsec



# Transport-layer security (TLS)

- widely deployed security protocol above the transport layer
  - supported by almost all browsers, web servers: https (port 443)
- provides:
  - **confidentiality**: via *symmetric encryption*
  - **integrity**: via *cryptographic hashing (MAC)*
  - **authentication**: via *public key cryptography*

} *all techniques we  
have studied!*


# Importance of TLS

- Originally designed for **secure e-commerce over http**, now used much more widely.
  - Retail customer access to online banking facilities.
  - Access to Gmail, Facebook, Yahoo, etc.
  - >70% of web traffic is now encrypted using TLS.
- TLS has become the de facto secure protocol of choice.
  - Used by hundreds of millions/billions of people and devices every day.




# HTTPS Adoption by Websites

Percentage of pages loaded over HTTPS in Chrome by platform



# TLS Threat Model

- End-to-end secure communications in
- the presence of a network attacker



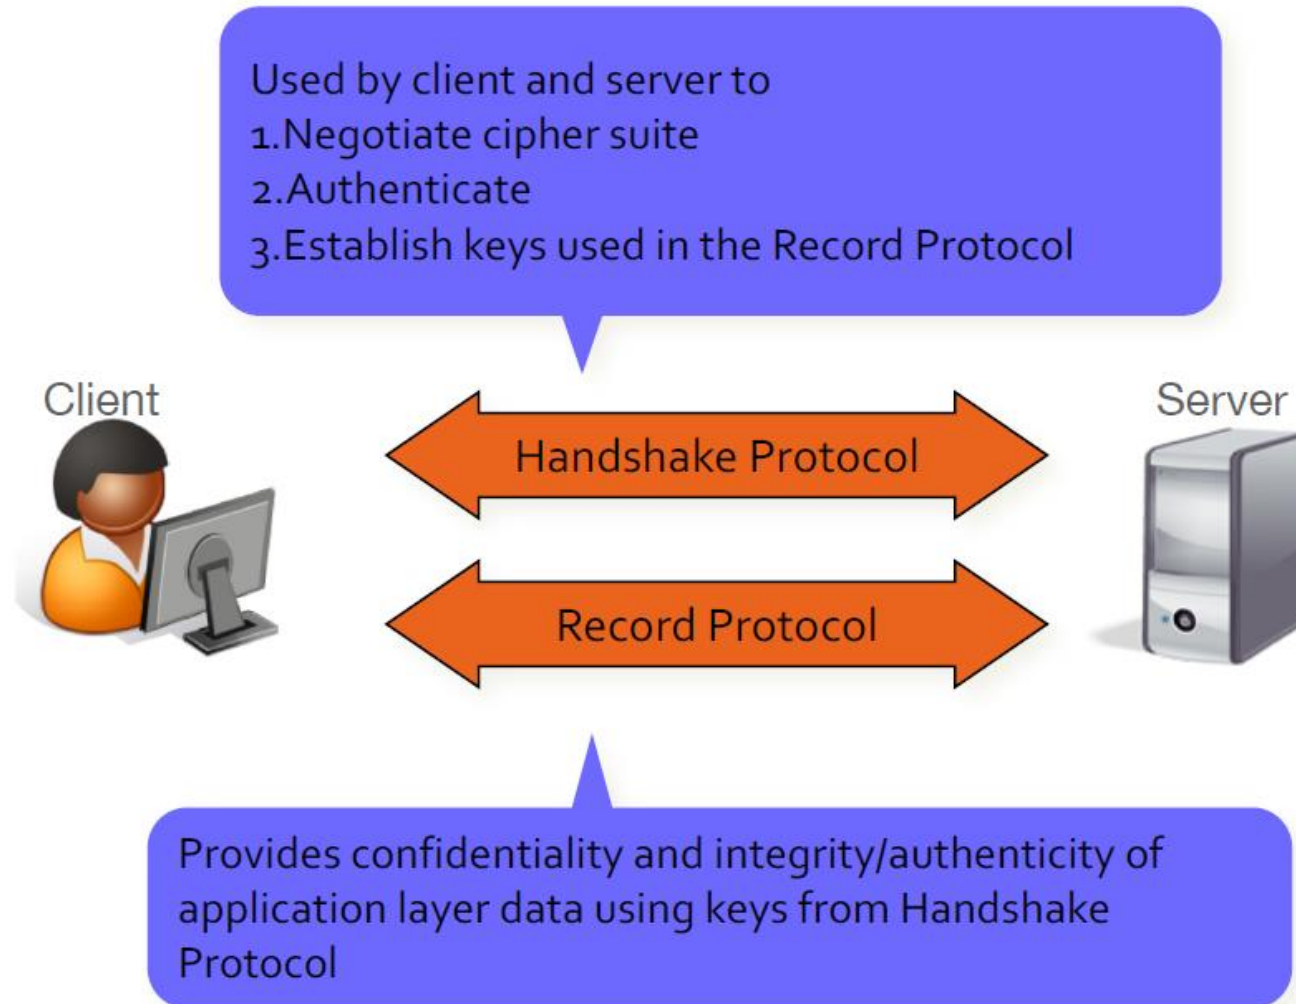
# SSL/TLS Overview

- SSL = Secure Sockets Layer.
  - Developed by Netscape in mid 1990s.
  - SSLv1 broken at birth.
  - SSLv2 seriously flawed in various ways.
  - SSLv3 now considered broken.
- TLS = Transport Layer Security.
  - IETF-standardised version of SSL.
  - TLS 1.0 in RFC 2246 (1999).
  - TLS 1.1 in RFC 4346 (2006).
  - TLS 1.2 in RFC 5246 (2008).
  - TLS 1.3 in RFC 8446 (2018).

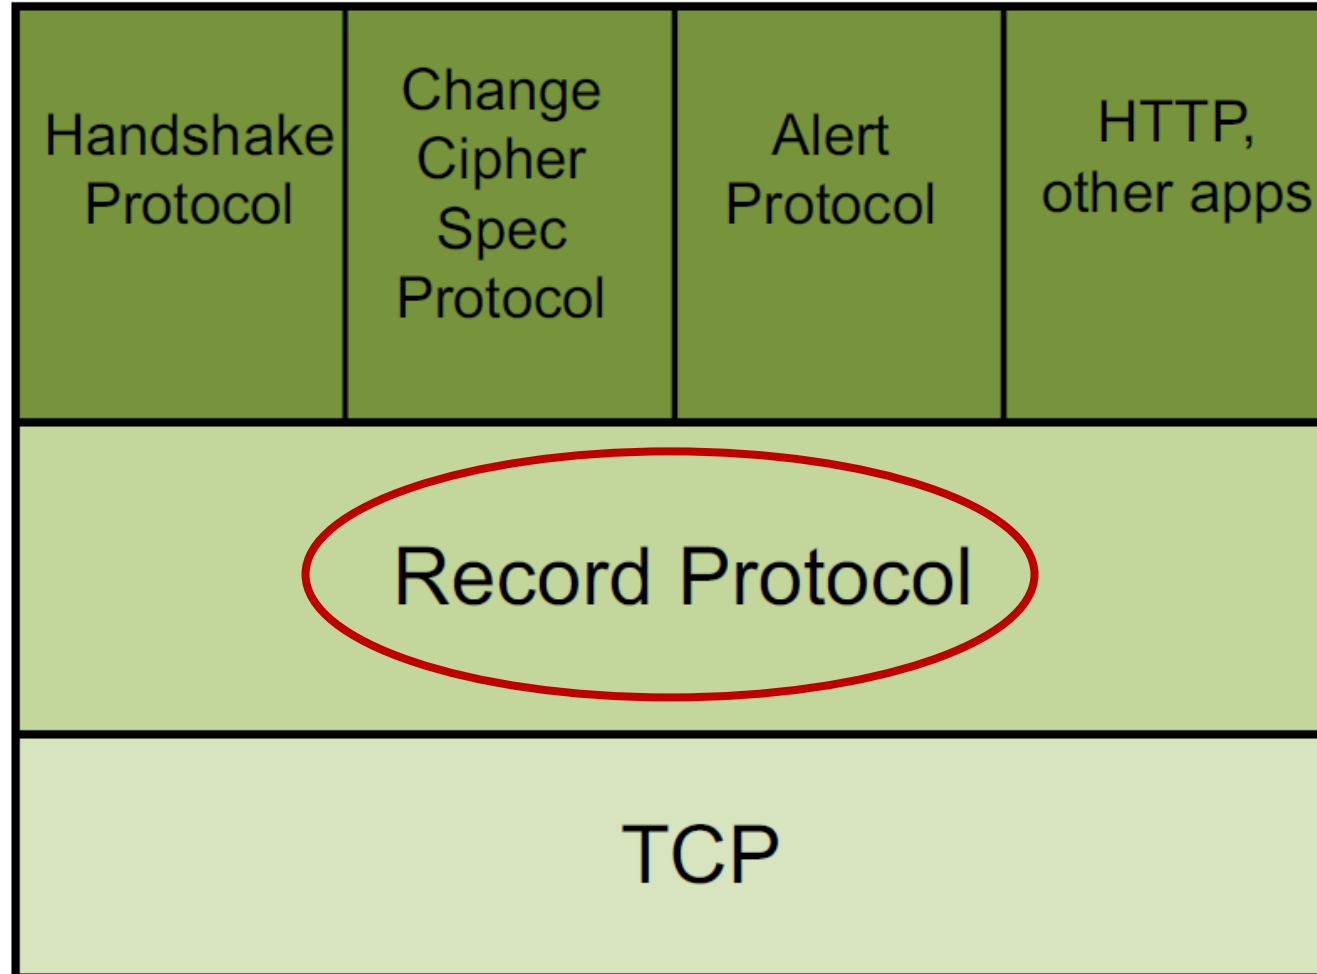
# Transport-layer security: what's needed?

- let's *build* a toy TLS protocol, *t-tls*, to see what's needed!
- we've seen the “pieces” already:
  - **handshake**: Alice, Bob use their certificates, private keys to authenticate each other, exchange or create shared secret
  - **key derivation**: Alice, Bob use shared secret to derive set of keys
  - **data transfer**: stream data transfer: data as a series of records
    - not just one-time transactions
  - **connection closure**: special messages to securely close connection

# Highly Simplified View of TLS



# TLS Protocol Architecture



# TLS: encrypting data

- recall: TCP provides data *byte stream* abstraction
- Q: can we encrypt data in-stream as written into TCP socket?
  - A: where would MAC go? If at end, no message integrity until all data received and connection closed!
  - solution: break stream in series of “records (fragments)”
    - each client-to-server record carries a MAC, created using  $M_c$
    - receiver can act on each record as it arrives
- t-tls record encrypted using symmetric key,  $K_c$ , passed to TCP:

$K_c$  ( 

<i>length</i>	<i>data</i>	<i>MAC</i>
---------------	-------------	------------

 )

# TLS: encrypting data (more)

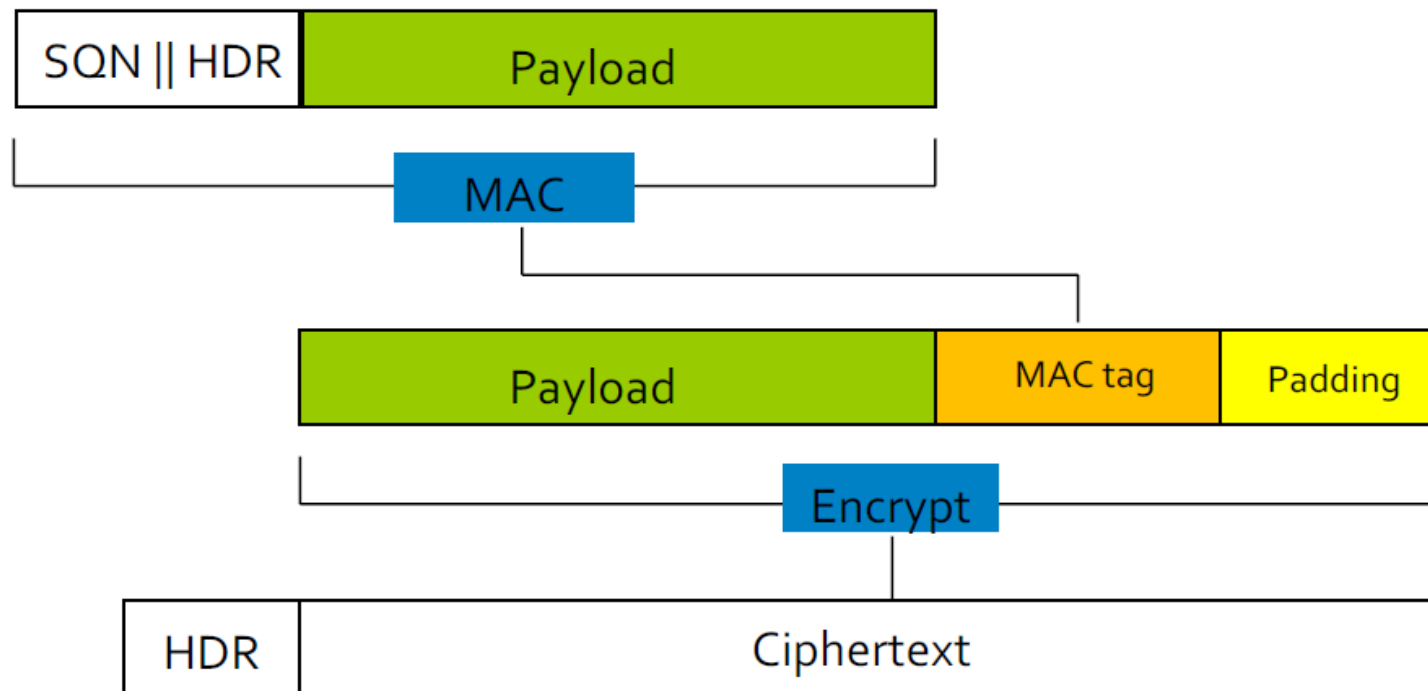
- possible attacks on data stream?
  - *re-ordering*: man-in middle intercepts TCP segments and reorders (manipulating sequence #s in unencrypted TCP header)
  - *replay*
- solutions:
  - use TLS sequence numbers (data, TLS-seq-# incorporated into MAC)
  - use nonce



# TLS Record Protocol

- **TLS Record** Protocol provides:
  - Data origin authentication, integrity using a MAC.
  - Confidentiality using a symmetric encryption algorithm.
  - Anti-replay using sequence numbers protected by the MAC.
  - Optional compression.

# TLS Record Protocol Operation



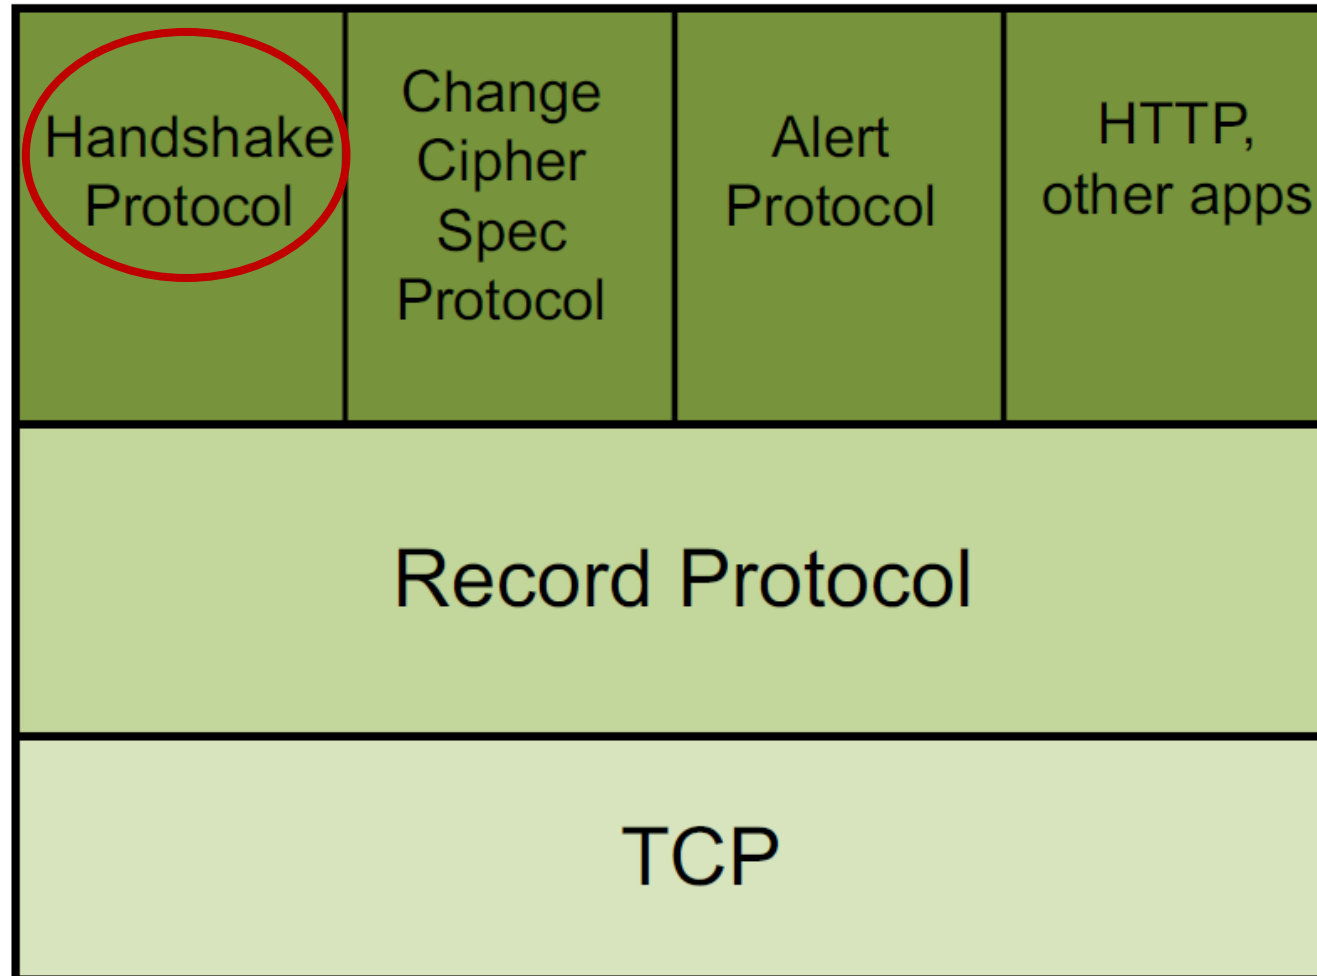
MAC

HMAC-MD5, HMAC-SHA1, HMAC-SHA256,...

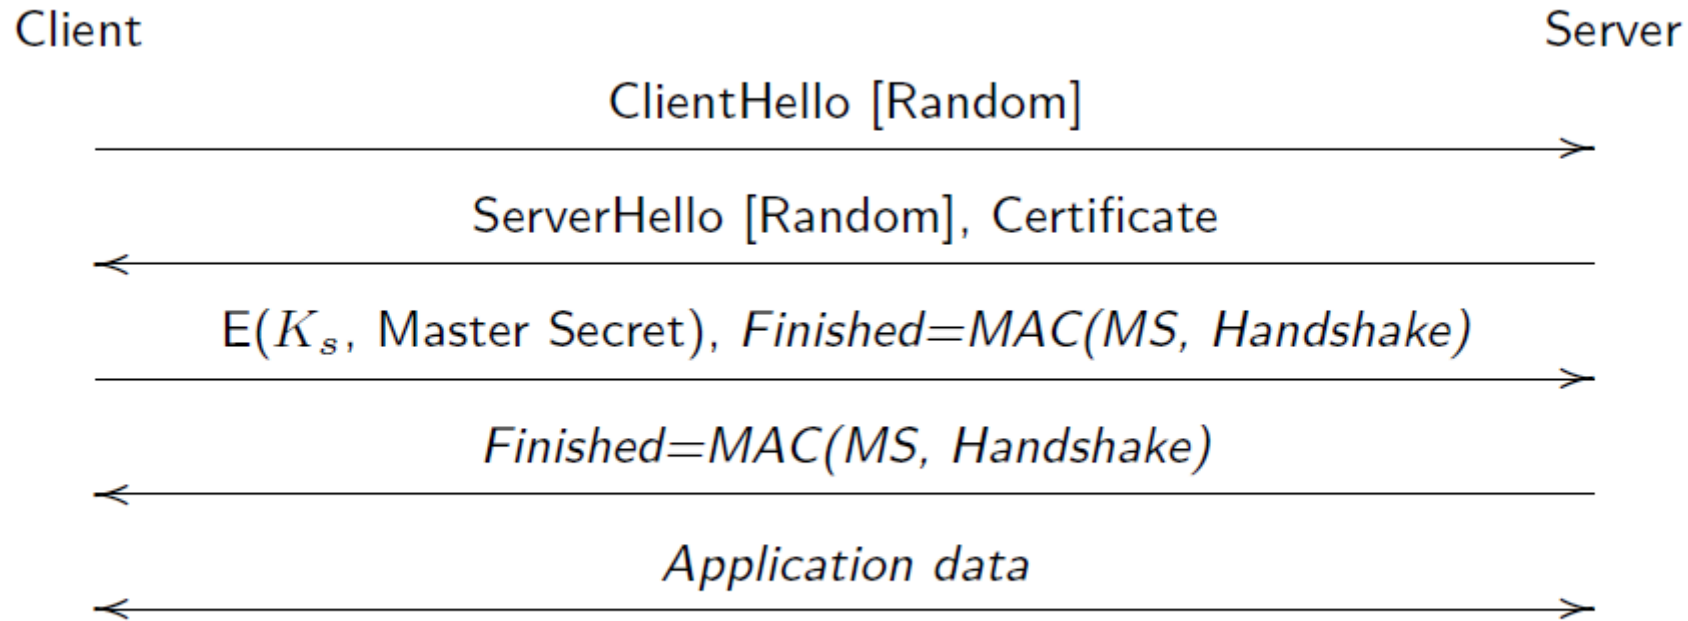
Encrypt

CBC-AES128, CBC-AES256, CBC-3DES, RC4-128,...

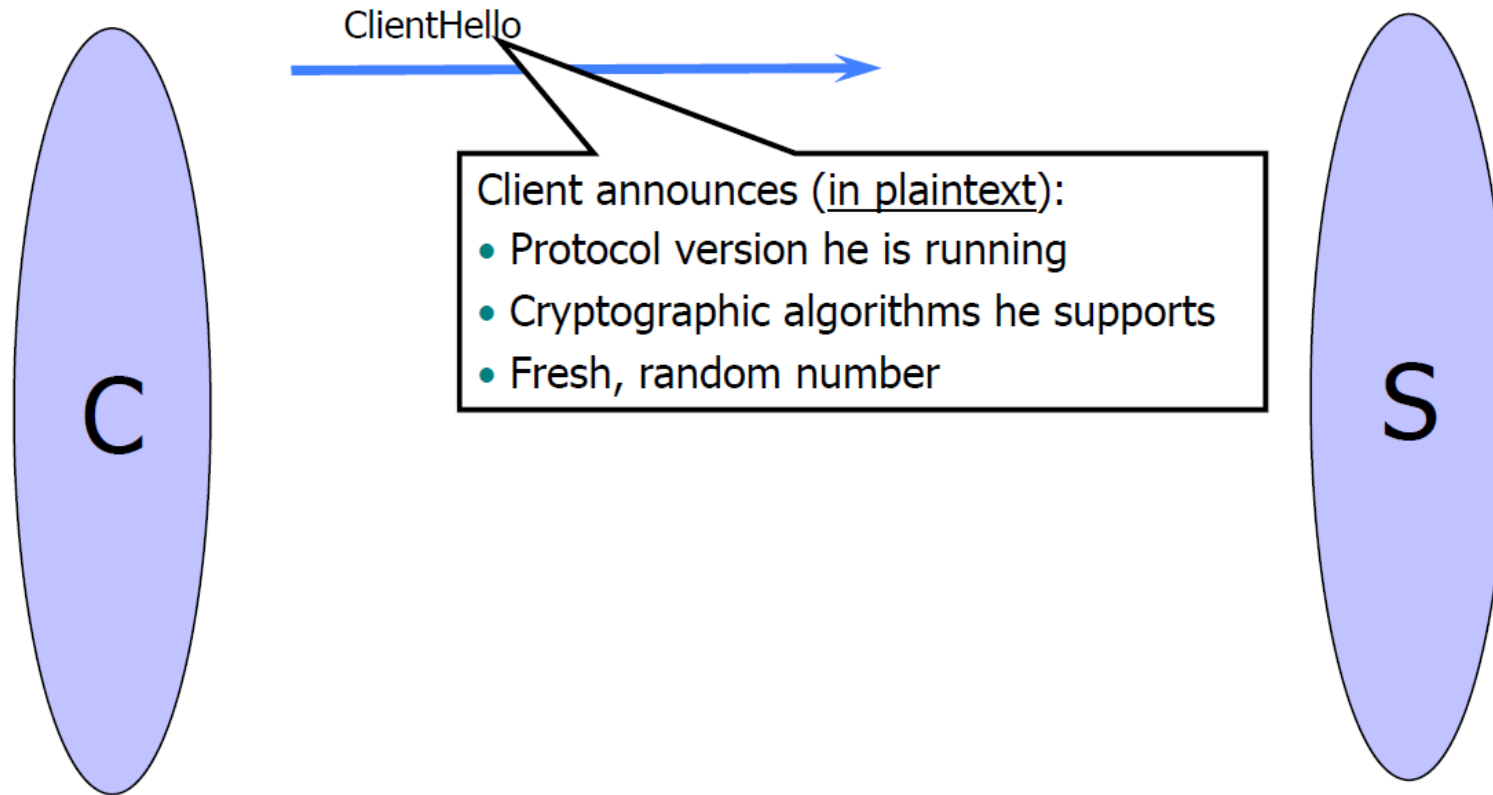
# TLS Protocol Architecture



# TLS 1.2: RSA Handshake Skeleton



# ClientHello



# ClientHello

```
struct {
```

```
    ProtocolVersion client_version;
```

Highest version of the protocol  
supported by the client

```
    Random random;
```

```
    SessionID session_id;
```

Session id (if the client wants to  
resume an old session)

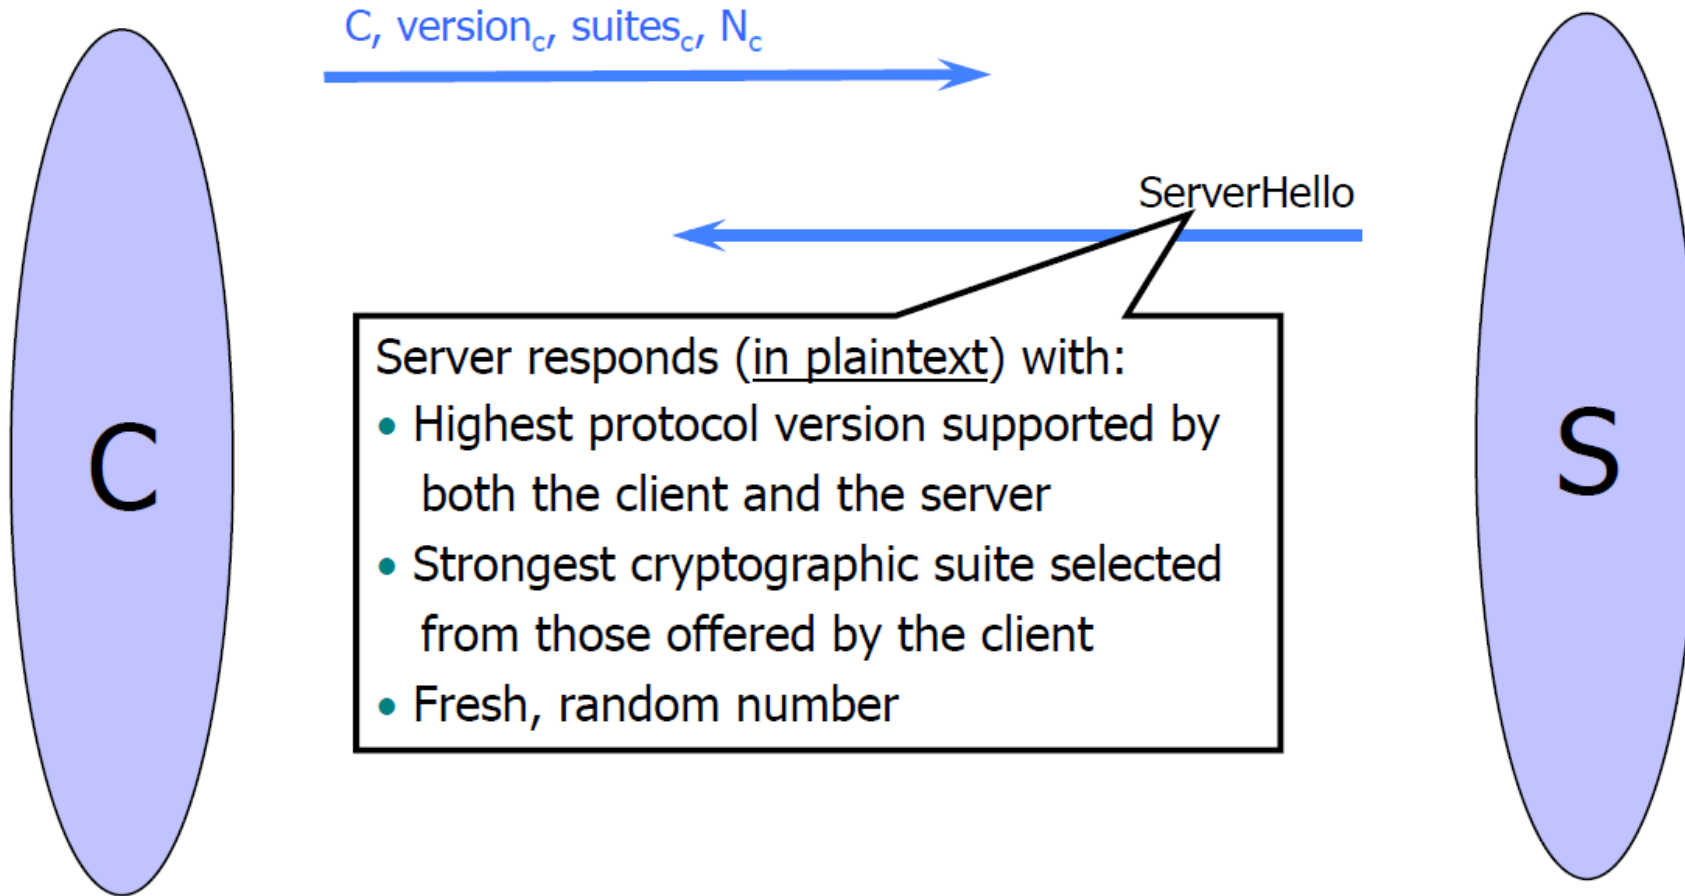
```
    CipherSuite cipher_suites;
```

Set of cryptographic algorithms  
supported by the client (e.g.,  
RSA or Diffie-Hellman)

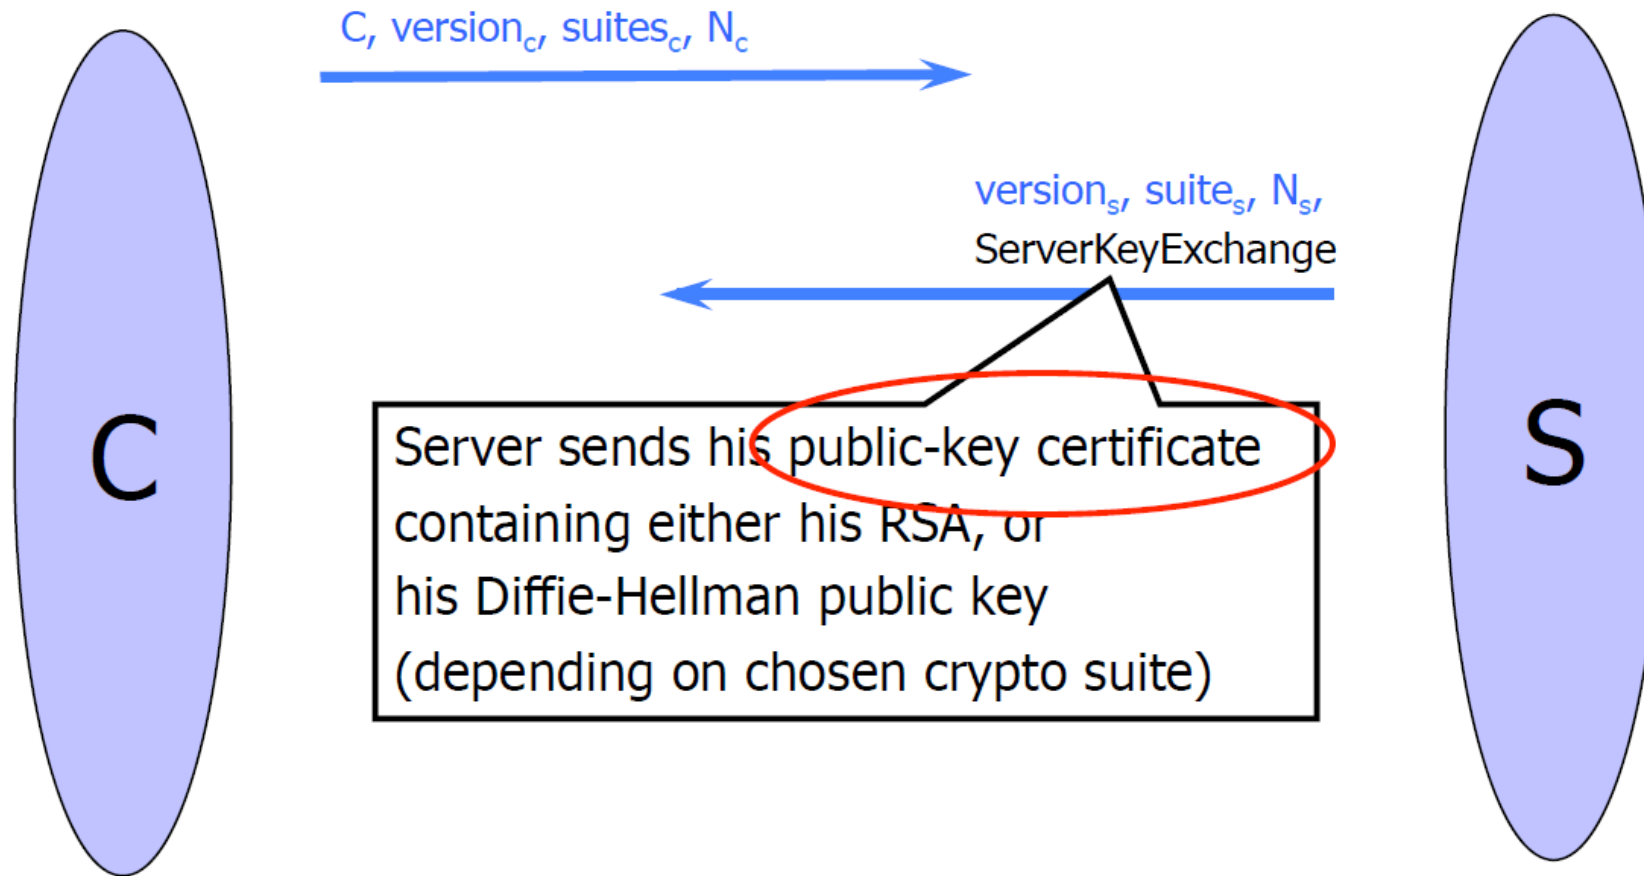
```
    CompressionMethod compression_methods;
```

```
} ClientHello
```

# ServerHello

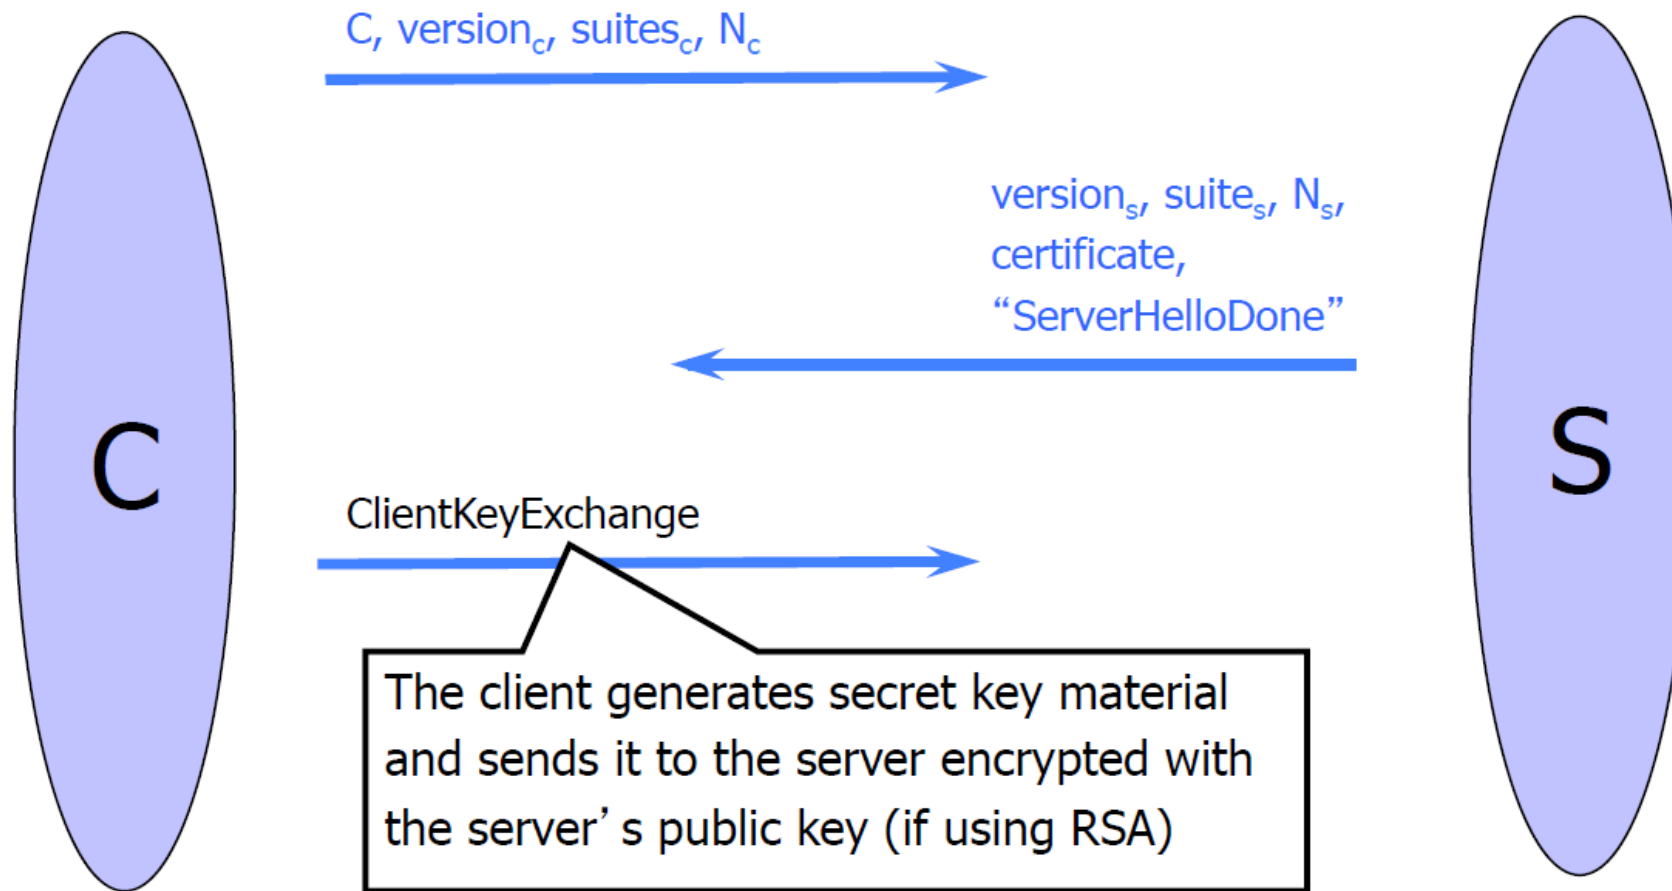


# ServerHello





# ClientKeyExchange



# ClientKeyExchange

```
struct {  
    select (KeyExchangeAlgorithm) {  
        case rsa: EncryptedPreMasterSecret;  
        case diffie_hellman: ClientDiffieHellmanPublic;  
    } exchange_keys  
} ClientKeyExchange
```

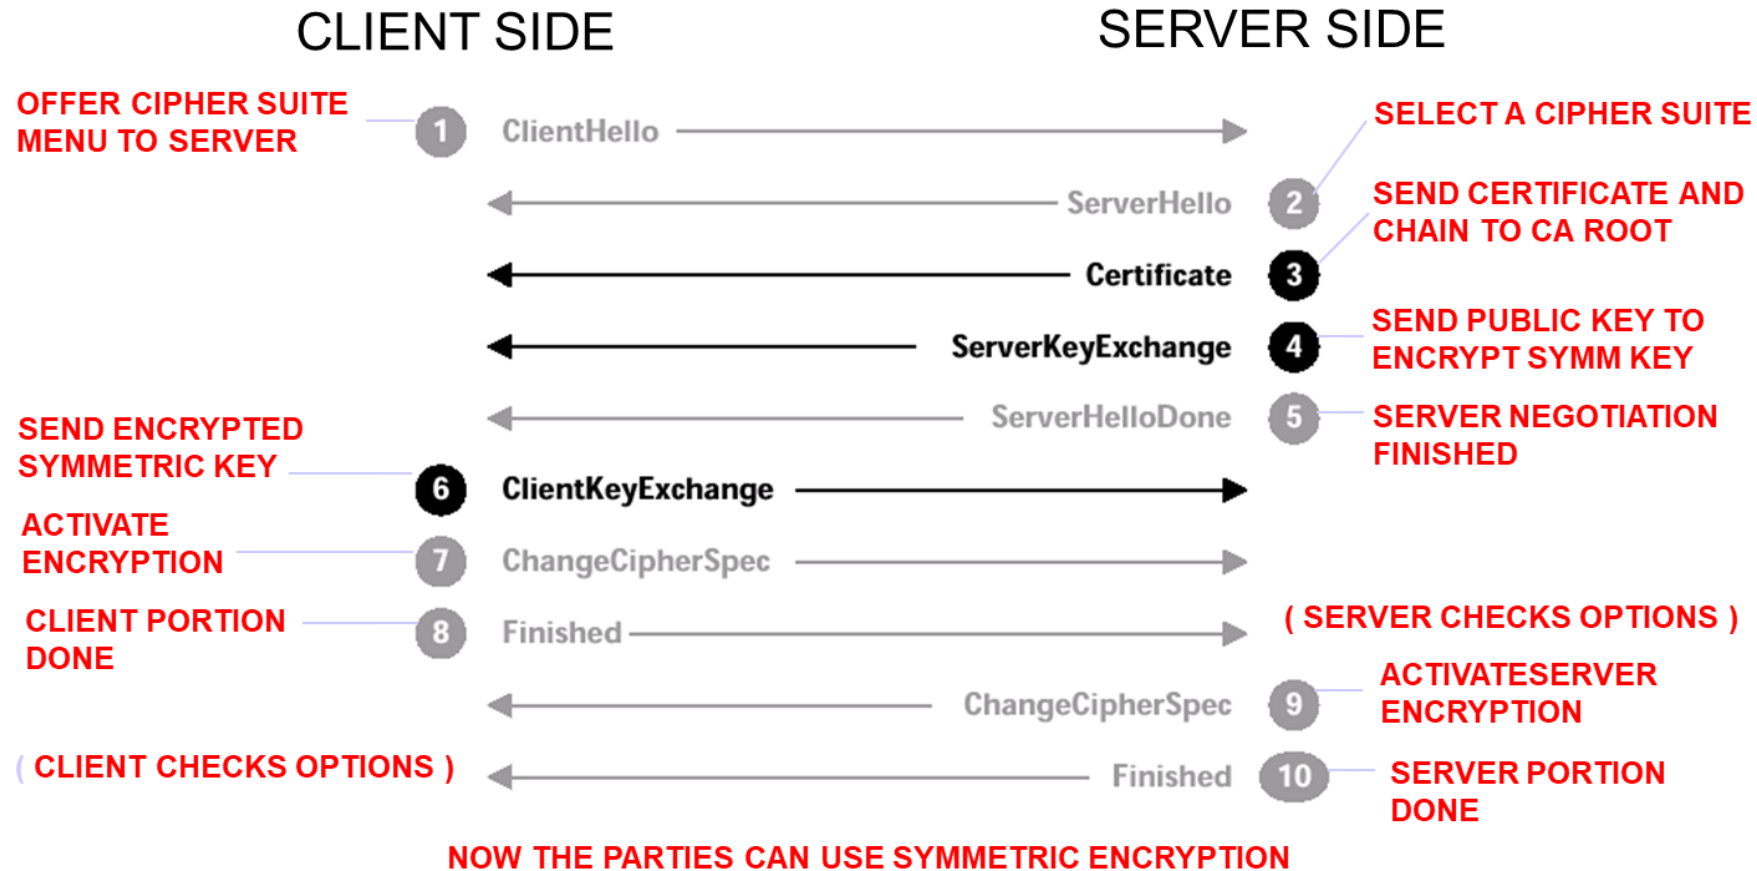
```
struct {  
    ProtocolVersion client_version;  
    opaque random[46];  
} PreMasterSecret
```

Where do random  
bits come from?



Random bits from which  
symmetric keys will be derived  
(by hashing them with nonces)

# TLS Handshake



SOURCE: THOMAS, *SSL AND TLS ESSENTIALS*

# TLS: cryptographic keys

- considered bad to use same key for more than one cryptographic function
  - different keys for message authentication code (MAC) and encryption
- four keys:
  - 🔑  $K_c$  : encryption key for data sent from client to server
  - 🔑  $M_c$  : MAC key for data sent from client to server
  - 🔑  $K_s$  : encryption key for data sent from server to client
  - 🔑  $M_s$  : MAC key for data sent from server to client
- keys derived from key derivation function (KDF)
  - takes master secret and (possibly) some additional random data to create new keys

# TLS 1.3 Objectives

- **Clean up**: Remove unused or unsafe features
- **Security**: Improve security by using modern security analysis techniques
- **Privacy**: Encrypt more of the protocol
- **Performance**: Our target is a 1-RTT handshake for naive clients; 0-RTT handshake for repeat connections
- **Continuity**: Maintain existing important use cases

# TLS: 1.3 cipher suite

- “cipher suite”: algorithms that can be used for key generation, encryption, MAC, digital signature
- TLS: 1.3 (2018): more limited cipher suite choice than TLS 1.2 (2008)
  - only 5 choices, rather than 37 choices
  - *requires* Diffie-Hellman (DH) for key exchange, rather than DH or RSA
  - combined encryption and authentication algorithm (“authenticated encryption”) for data rather than serial encryption, authentication
    - 4 based on AES
  - HMAC uses SHA (256 or 284) cryptographic hash function

# References

- Kurose, James F., and Keith W. Ross. "Computer networking: A top-down approach edition." Addison Wesley (2007), chapter 8.
- Cryptography and Network Security: Principles and Practice, William Stallings, Pearson, 2022.
- COMS W4182 — Computer Security II (Spring '22), Prof. Steven M. Bellovin, Columbia University.
- CS 5435 - Fall 2022 “Security and Privacy Concepts in the Wild” Vitaly Shmatikov, Cornell University.
- Introducing TLS, Information Security Group, Kenny Paterson.