



Advanced Network Security


Firewalls and IDS

Amir Mahdi Sadeghzadeh, Ph.D.

Firewalls

firewall

isolates organization's internal network from larger Internet, allowing some packets to pass, blocking others



Limitations of firewalls, gateways


- **Interfere** with some networked applications
- Don't solve many real problems
 - **Buggy software** (think buffer overflow exploits)
 - **Bad protocol design** (think WEP in 802.11b)
- Generally don't prevent **denial of service**
- Don't prevent **insider attacks**
- Increasing complexity and potential for **misconfiguration**

Intrusion detection systems

- packet filtering:
 - operates on TCP/IP headers only
 - no correlation check among sessions
- IDS: intrusion detection system
 - deep packet inspection: look at packet contents (e.g., check character strings in packet against database of known virus, attack strings)
 - examine correlation among multiple packets
 - port scanning
 - network mapping
 - DoS attack

Intrusion detection systems

multiple IDSs: different types of checking at different locations



Fighting intrusion

- Prevention: isolate from network, strict authentication measures, encryption
- Preemption:
 - “do to others before they do to you”
- Deterrence: dire warnings,
 - “we have a bomb too.”
 - Mutual Assured Destruction (MAD)
- Deflection: diversionary techniques to lure away
- Detection
- Counter attacks

Defense in Depth

- More generically, most single defenses can fail
- We always need defense in depth – multiple layers, of different designs and philosophies
- One such layer: Intrusion Detection Systems

What is IDS?

- An Intrusion Detection System (IDS) is a system that attempts to **identify intrusions**.
- Intrusion detection is the process **of identifying and responding to malicious activity** targeted at **computing and networking resources**.
- The goal of IDS is to **detect fingerprints** of malicious activity.

Examples of IDS in daily life

- Car Alarms
- House Alarms
- Surveillance Systems
- Spy Satellites, and spy planes

Elements of Intrusion Detection

- Primary assumptions:
 - System activities are **observable**
 - **Normal and intrusive activities have distinct evidence**
- Components of intrusion detection systems:
 - From an algorithmic perspective:
 - **Features** - capture intrusion evidence from audit data
 - **Models** - piece evidence together; infer attack
 - From a system architecture perspective:
 - Audit data processor, knowledge base, decision engine, alarm generation and responses

Where Are IDS Deployed?

■ Host-based

- Monitor activity on a single host
- Advantage: **better visibility** into behavior of individual **applications** running on the host and network traffic

■ Network-based (NIDS)

- Often placed on a **router or firewall**
- Monitor traffic, examine **packet headers and payloads**
- Advantage: single NIDS can protect many hosts and **look for global patterns**

Host-Based IDSs

- Using OS auditing mechanisms
- BSM on Solaris: logs all direct or indirect events generated by a user
- Strace (the linux syscall tracer) for system calls made by a program
- Monitoring user activities
 - E.G., Analyze shell commands
- Monitoring execution of system programs
 - E.G., Analyze system calls made by sendmail

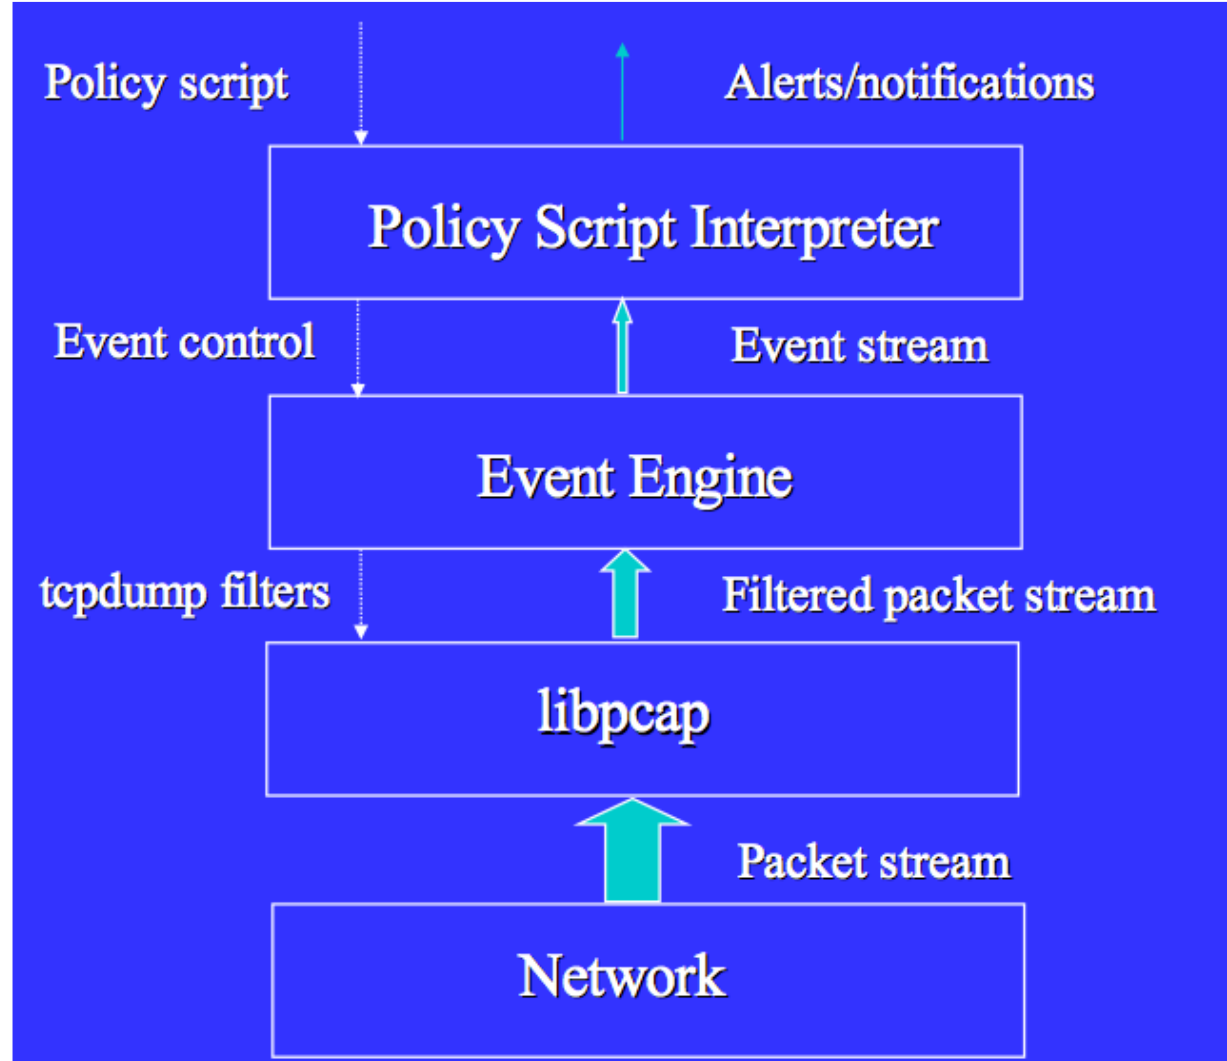
Basic Audit Modules (Hosts)

- eventLog - Uses the windows Event Logging system to track entries into all three of the windows event logs: System, Security, Application
- netstat - Uses the information from the program netstat to provide information about network usage on the machine
- health - Runs the program health to give current information about the system (CPU usage, mem usage, swap usage)
- ps - Uses information from the /proc virtual file system as a data source

Network IDSs


- Deploying **sensors** at **strategic locations**
 - E.G., Packet sniffing via tcpdump at routers
- Inspecting network traffic
 - Watch for **violations of protocols** and **unusual connection patterns**
- Monitoring user activities
 - Look into the **data portions of the packets for malicious command** sequences
- Problem
 - May be easily **defeated by encryption**
 - Data portions and some header information can be encrypted
 - Data volume

Architecture of Network IDS



Firewall Versus Network IDS

- Firewall
 - Active filtering
 - Fail-close
- Network IDS
 - Passive monitoring
 - Fail-open
- IPS (IDS + Firewall)



Requirements of Network IDS

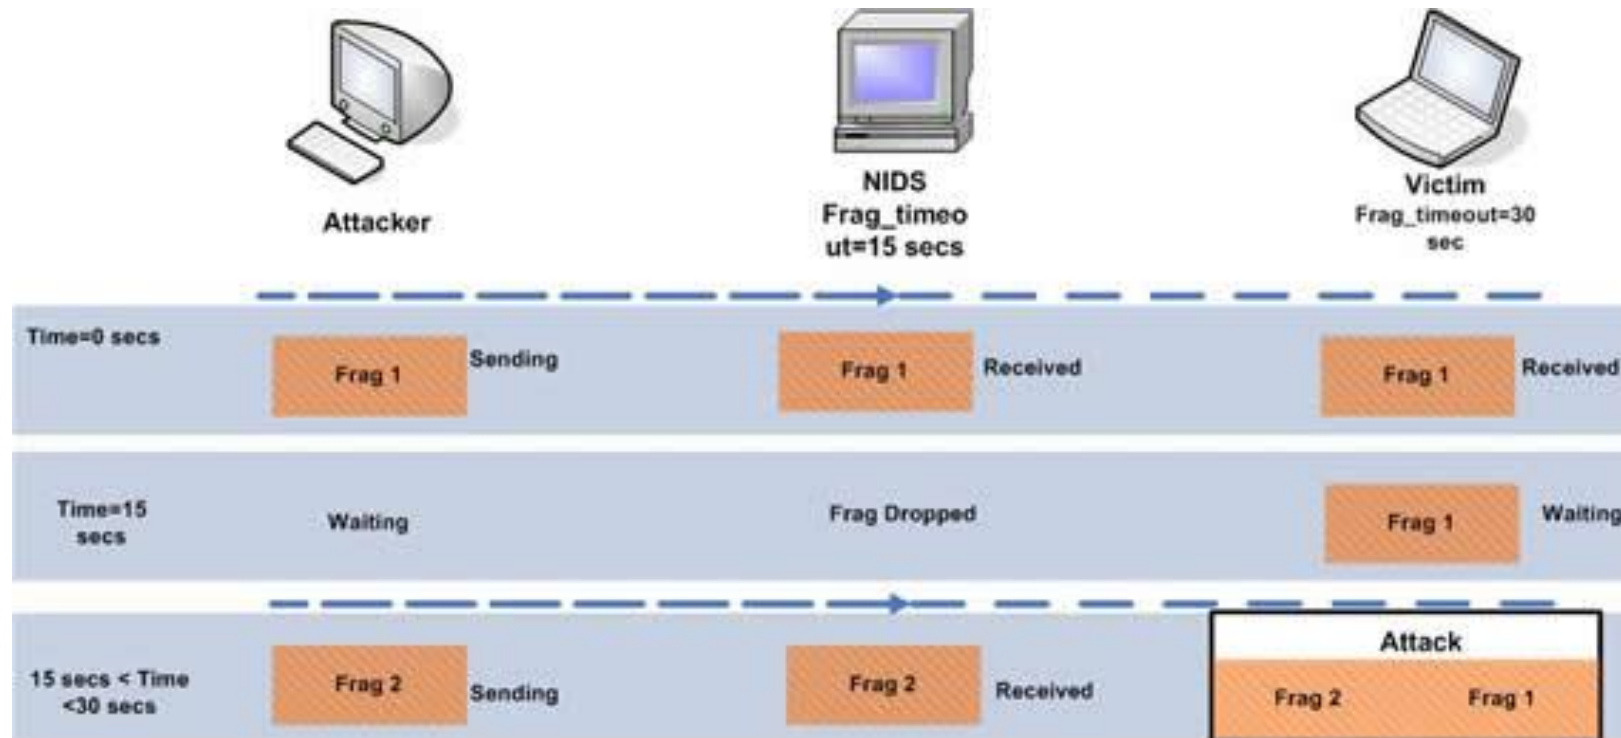
- High-speed, large volume monitoring
 - No packet filter drops
 - Why is it hard?
- Real-time notification
- Broad detection coverage
 - Precision, Recall, F-score
- Economy in resource usage
- Resilience to stress
- Resilience to attacks upon the IDS itself!

Eluding Network IDS

- What the **IDS sees may not be what the end system gets.**
 - Insertion and evasion attacks.
 - IDS needs to perform **full reassembly of packets.**
- But there are still **ambiguities in protocols and operating systems:**
 - E.G. TTL, fragments.
 - Need to “normalize” the packets.

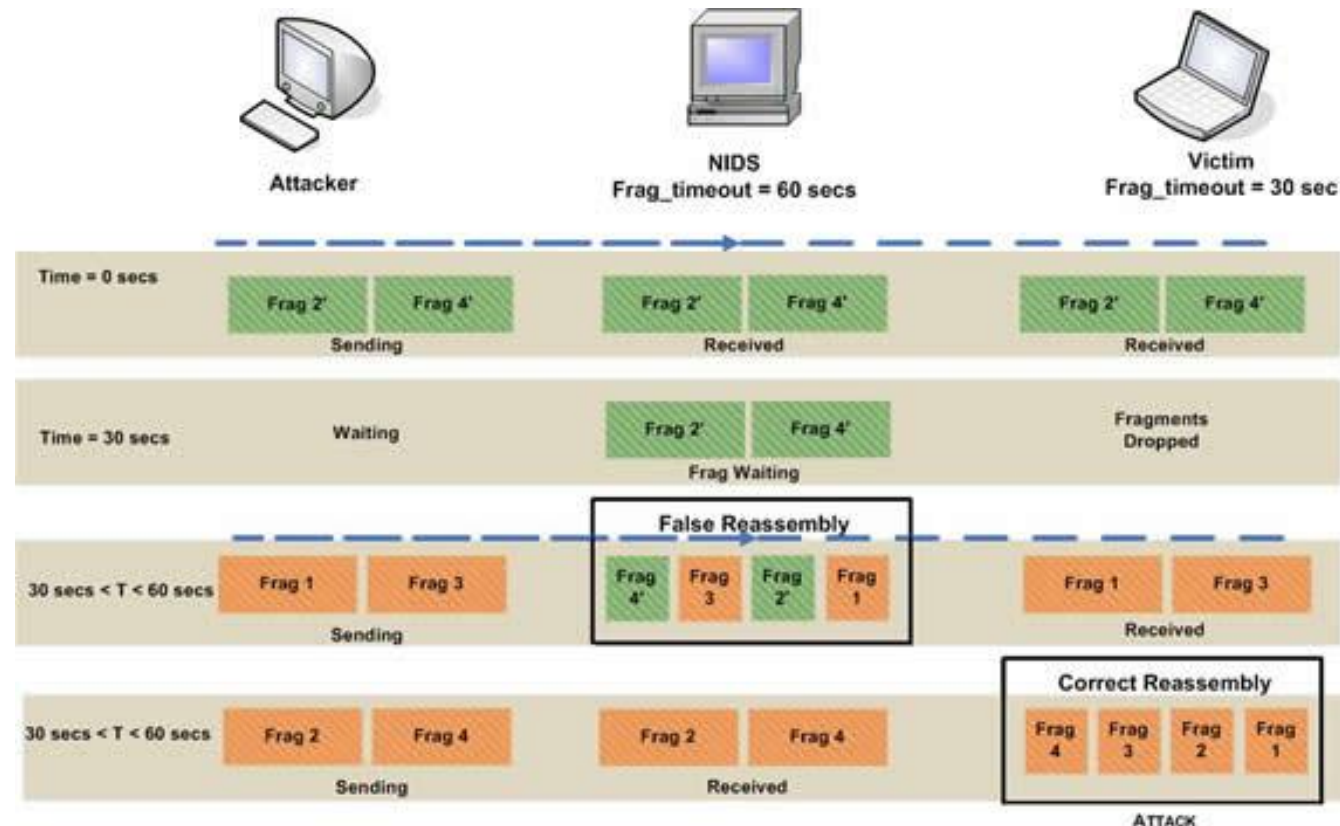
Insertion Attack

- IDS fragmentation reassembly timeout is 15 seconds
- The system is monitoring some Linux hosts which have a default fragmentation reassembly timeout of 30 seconds.

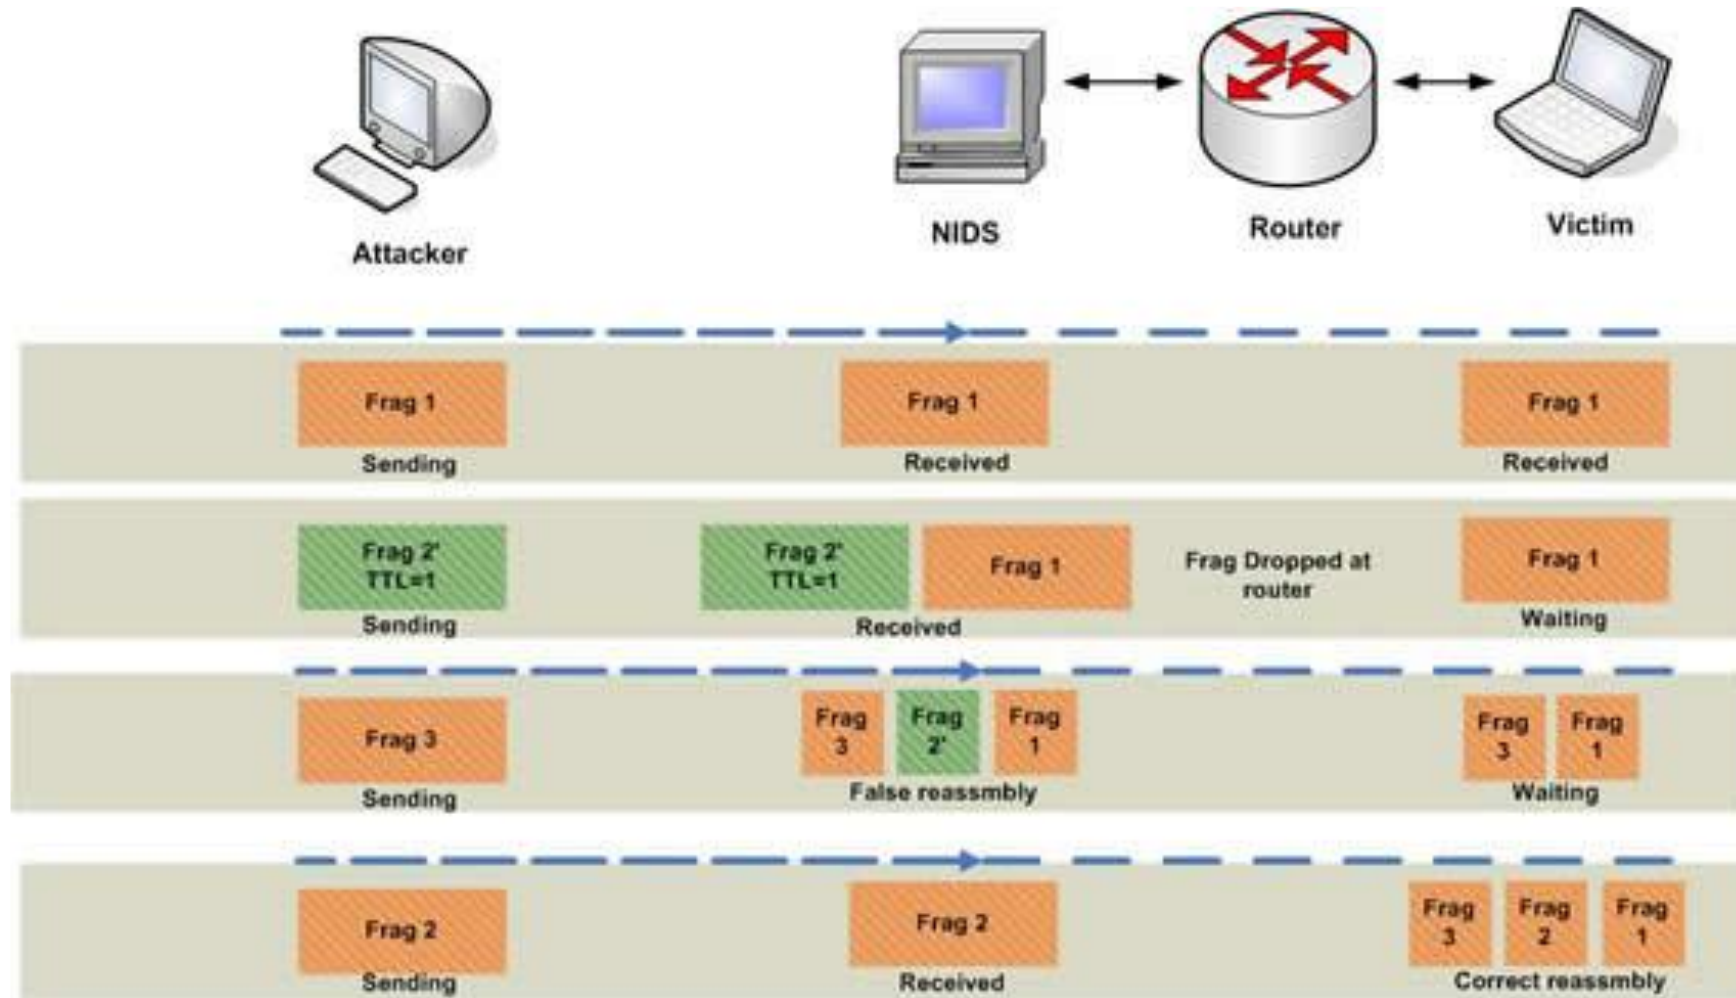


Insertion Attack

- By default, Snort has a fragment reassembly timeout of 60 seconds.
- Compare that to Linux/FreeBSD where it is 30 seconds.



Insertion Attack



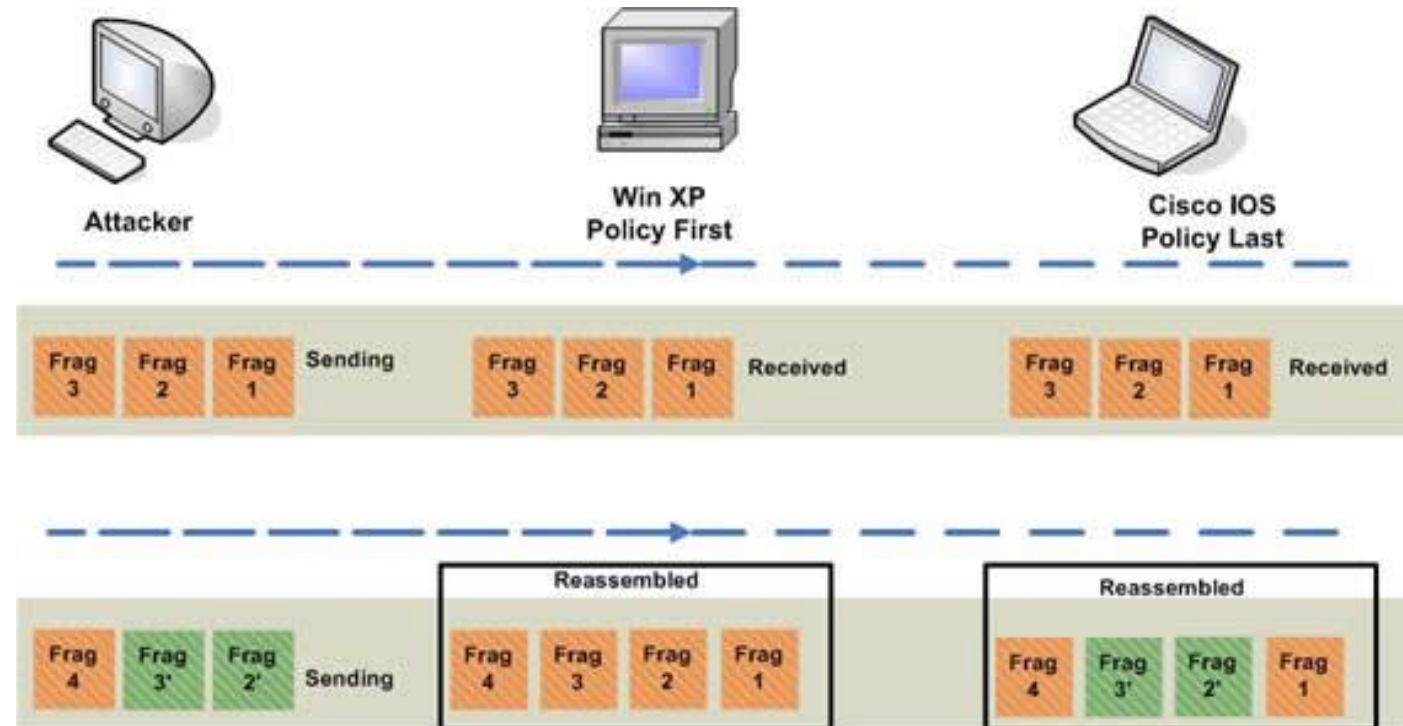
Insertion Attack

■ First

- This is where the operating System favors the original fragments with a given offset. For example, Windows 95/98/NT4/ME/W2K/XP/2003.

■ Last.

- This is where the operating System favors the subsequent fragments with a given offset. For example, Cisco IOS.



Snort evasion countermeasures


- Snort is the most popular NIDS.
 - The frag3 preprocessor is a target-based IP defragmentation module for Snort.
 - Allowing a user to identify the **fragmentation reassembly method and the corresponding fragment timeout value** that is applied to a particular destination IP address or subnet.
 - preprocessor frag3_engine: -policy bsd -bind_to 192.168.1.0/24 -timeout 30 -min_ttl 2

DoS Attacks on Network IDS

- Resource exhaustion
 - CPU resources
 - Memory
 - Network bandwidth
- Abusing reactive IDS
 - False positives

Hybrid NIDS and HIDS

- Sensors
 - Trust issue



Hybrid NIDS and HIDS

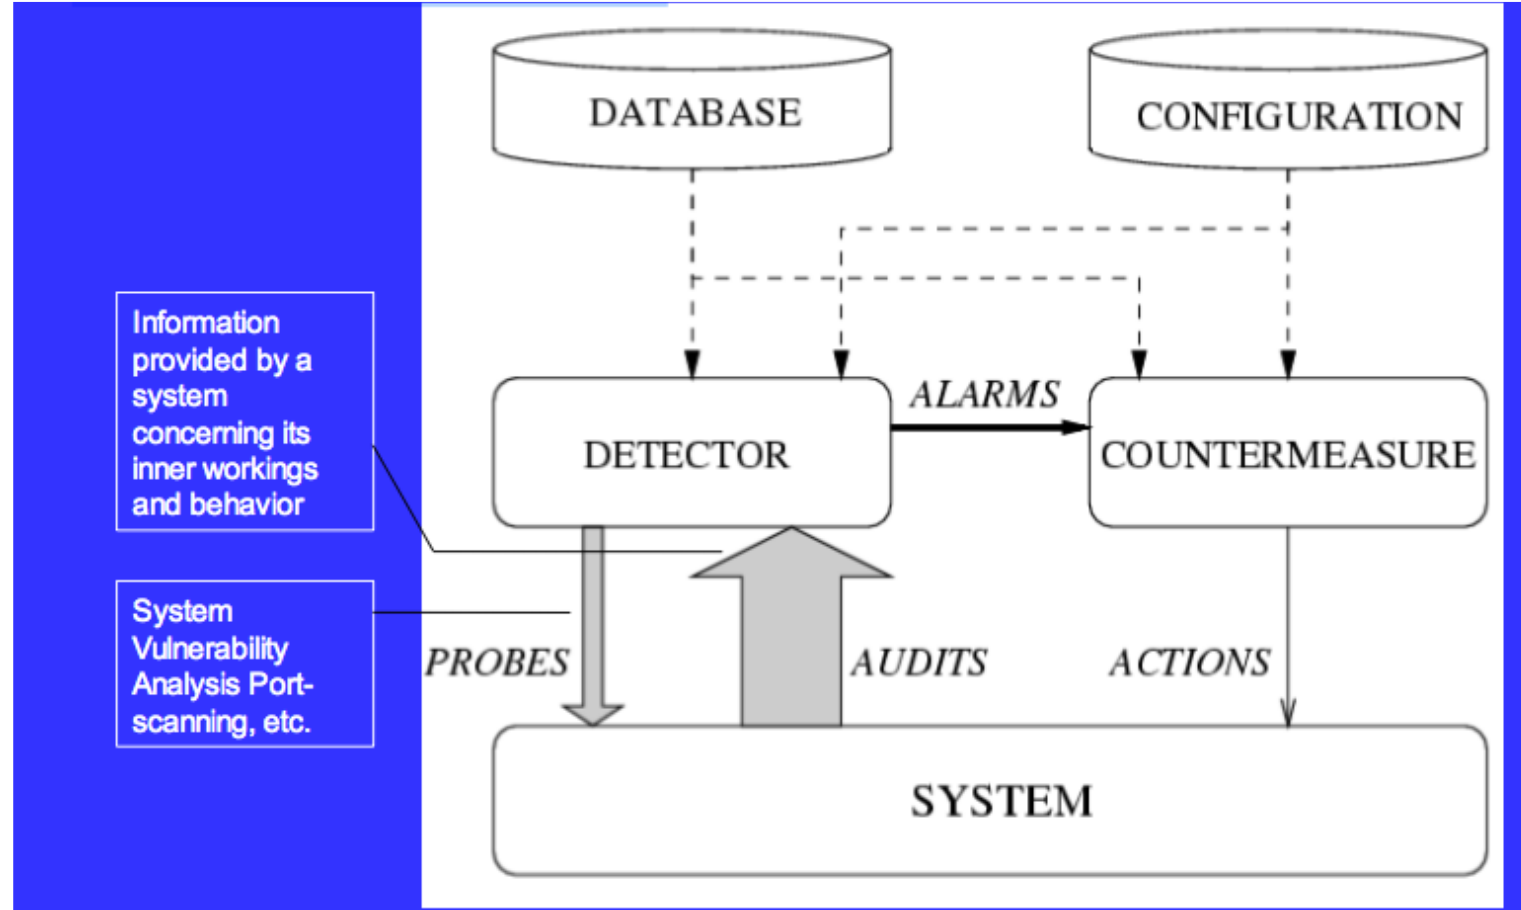
- Correlate information from multiple sources
- How do you trust your sources?

Taxonomy of IDS's

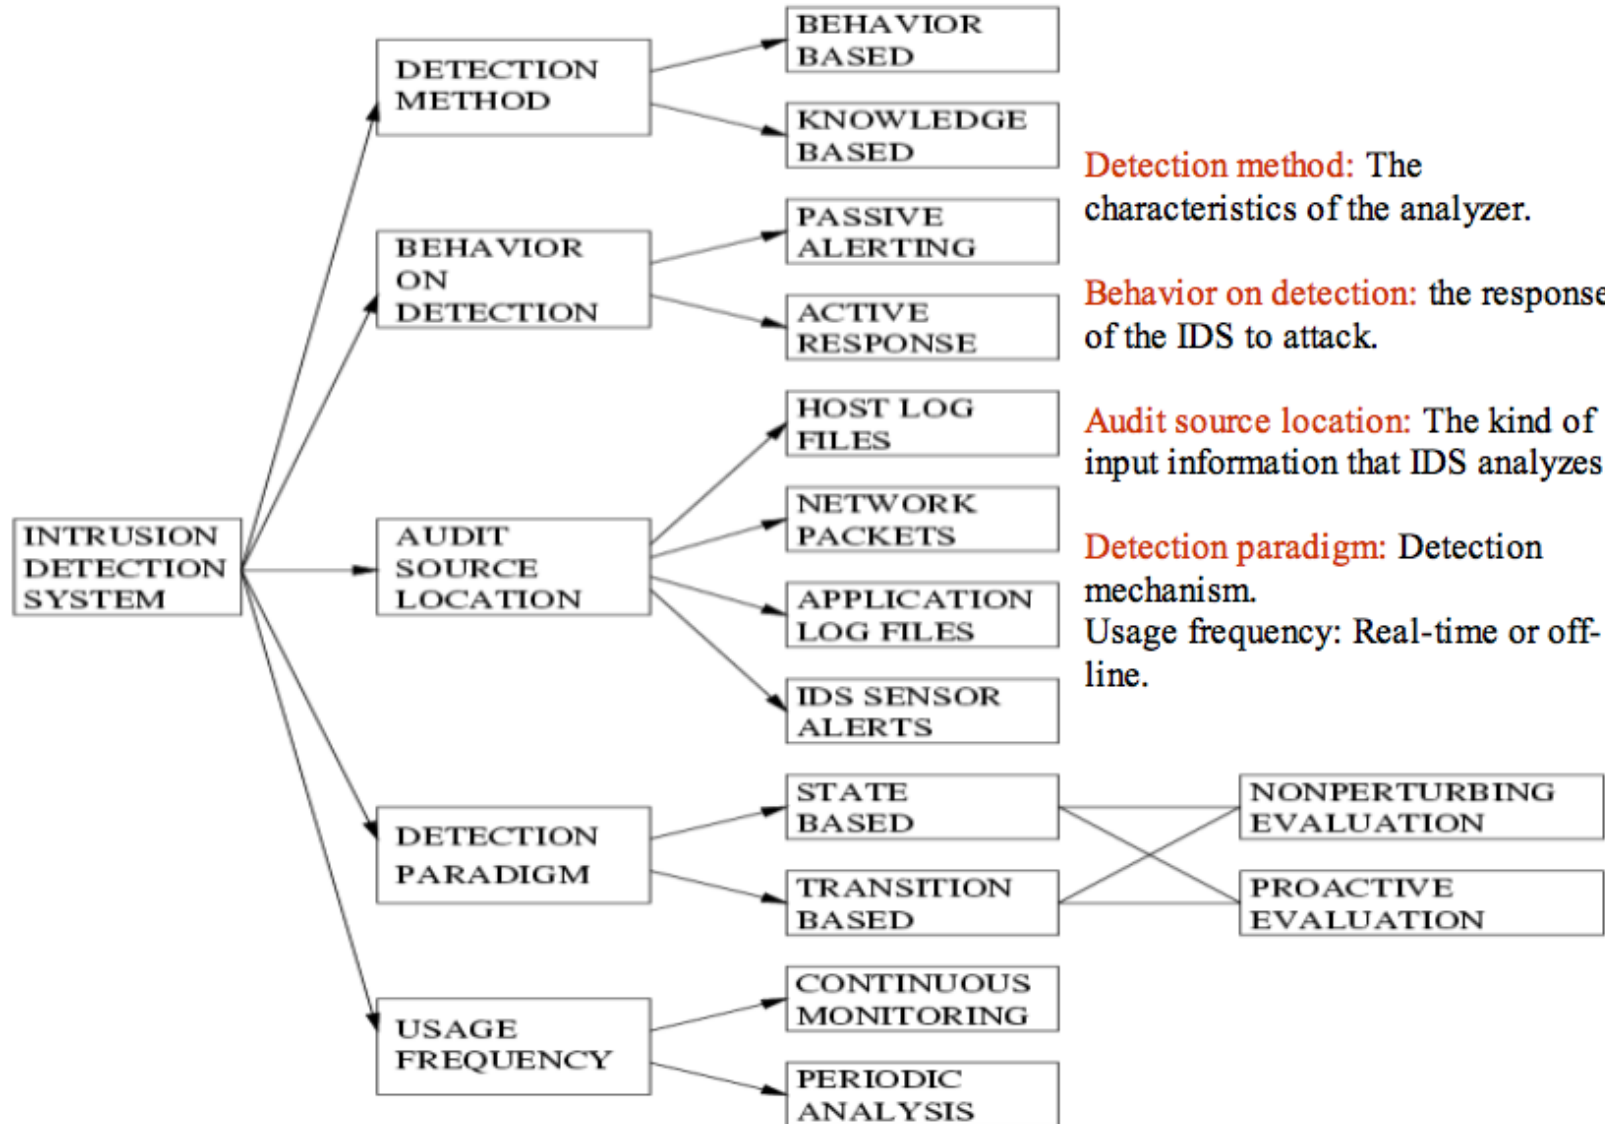
Intrusion Detection Approaches

- Modeling
 - Features: evidences extracted from audit data
 - Analysis approach: piecing the evidences together
 - Misuse detection (a.k.a. signature-based)
 - Anomaly detection (a.k.a. statistical-based)
- Deployment: Network-based or Host-based
- Development and maintenance
 - Hand-coding of “expert knowledge”
 - Learning-based on audit data

A Generic IDS



Characteristics of IDS



Detection Paradigm

- State-based versus transition-based IDS
 - **State-based**: Identifies intrusions on the states
 - **Transition-based**: Watches events that trigger transition from one state to another
- Non-perturbing versus pro-active analysis of state or transition
 - **Non-perturbing**: Acquire information transparently
 - **Pro-active**: Analysis by explicitly triggering events

IDS: Time aspect

■ Real-time IDS

- Analyzes the data while the sessions are in progress
- Raises an alarm immediately when the attack is detected

■ Off-line IDS

- Analyzes the data after the information has been already collected
- Useful for understanding the attackers' behavior

Knowledge-based IDS

- Good accuracy, bad completeness
 - Drawback
 - need regular update of knowledge
 - Difficulty of gathering the information
 - Maintenance of the knowledge is a time-consuming task
- Knowledge-based IDS
 - Misuse Detection
 - Specification-based Detection

Misuse Detection

- The system is equipped with a number of **attack descriptions** (“**signature**”).
 - Then matched against the audit data to detect attacks.
- **Signature**
 - Sequences of system calls, patterns of network traffic, etc
- **Pro**: less **false positives** (But there still some!)
- **Con**: cannot detect **novel attacks**, need to update the signatures often.
- Approaches: pattern matching, security rule specification.

Misuse Detection (Signature-Based)

- Set of **rules** defining a behavioral signature likely to be associated with attack of a certain type
 - Example: **buffer overflow**
 - A setuid program spawns a shell with certain arguments
 - A network packet has lots of NOPs in it
 - A very long argument to a string function
 - Example: **SYN flooding** (denial of service)
 - Large number of SYN packets without ACKs coming back
- Attack signatures are usually very specific and **may miss variants of known attacks**
 - Why not make signatures more general?

Extracting Misuse Signatures

- Use **invariant** characteristics of known attacks
 - **Bodies of known viruses and worms, port numbers** of applications with known buffer overflows, RET addresses of stack overflow exploits
 - Hard to handle malware mutations
 - **Metamorphic** viruses: each copy has a different body
- **Challenge**: fast, automatic extraction of signatures of new attacks
- **Honeypots** are useful for signature extraction
 - Try to attract malicious activity, be an early target

Specification-based Detection

- Manually develop **specifications that capture legitimate** (not only previous seen) system behavior (**all good states**).
- **Any deviation** from it is an **attack**

- **Pro**: can avoid **false-positive** since the specification can capture all legitimate behavior.
- **Con**: **hard to develop** a complete and detailed specification, and error-prone.
- Approach: state machines

Today's IT Security Tools

- We make lists of bad behavior
 - Virus definitions
 - SPAM filters and blacklists
 - IDS signatures
 - Policies
- We distribute the lists to applications and detection systems
- They **flag behavior that fits the pattern**
- The system is **about to collapse**
 - Delays
 - Administrative Overhead
 - False positives

References

- Kurose, James F., and Keith W. Ross. "Computer networking: A top-down approach edition." Addison Wesley (2007), chapter 8.
- Steven Bellovin, COMS W4180, Columbia University, 2006
- Mehdi Kharrazi, CE40-817, Sharif University of Technology, 2015
- Vitaly Shmatikov, CS 361S, UT Austin, 2014