



Advanced Network Security

Worms and Viruses

Amir Mahdi Sadeghzadeh, Ph.D.

Worms and Viruses

- What are they?
- How do they spread?
- What can be done about them?

Viruses

- “Infected” program (or floppy)
- When program is executed, it performs its normal function
- It also infects some other programs
- It may carry an extra “payload” that performs other functions

Worms

- Similar to viruses, but they spread between machines
- Some are fully automatic; some require manual intervention to spread
- Some exploit bugs; others use social engineering

Classic Worms

Early Worms

- IBM Christmas Card “Virus”, December 1987
- Morris Internet Worm, November 1988
- Most worms since then have emulated one or both of those

Christmas Card Virus

- Infected EARN, BITNET, and IBM's VNET
- (Old, pre-TCP/IP network for IBM mainframes)
- Spread by **social engineering**

What Users Saw

- The program arrives in an email with the subject line "Let this exec run and enjoy yourself!".
- A comment inside the Christmas Tree source code contains the comment:
 - Browsing this file is no fun at all. Just type Christmas.

[illegible]

What Happened

- A file transfer mechanism (not quite email, though it could have been) delivered a short script to users
- It was written in REXX, a shell script-like language for IBM's VM/CMS system
- When **executed**
 - **Christmas Tree displays** an ASCII Christmas tree.
 - It reads the **files NAMES and NETLOG**, files **containing the addresses of communication partners**, and **mails** itself to every email address in them.
 - It transmitted a **copy of itself to any usernames it found**
- People **trusted** it, because it was coming from **a regular correspondent**. . .

Essential Elements

- Self-replicating executable
- Apparently from a trusted source
- Request that the recipient execute the program
- Using the NETLOG file to find new victims

These characterize most current email worms

The Damage

- The worm itself **wasn't malicious**
- However, it had **exponential growth patterns**
- It **clogged servers**, communication **paths**, **spool directories**, etc.
- In other words, it was an **unintentional denial of service** attack

The Internet Worm

- Also known as the **Morris worm**
- Got much more **mainstream publicity**
- Estimated to have taken out **6000 hosts — 10% of the Internet**
- Arguably, the first time the **Internet made the evening news**

Characteristics

- Much more **sophisticated**
- Exploited **buggy code** — spread **without human intervention**
- Exploited **trust patterns** among computers
- **Multiple attack vectors**
- **Multiple architectures** (Vax and Sun 3)
- Intended to **demonstrate the insecurity** of the **Internet**. . .

Attack Vectors

- Back door in sendmail
- Buffer overflow in fingerd
- Password-guessing
- Pre-authenticated login via rsh

Sendmail Back Door

- The author of sendmail wanted continued access to the production version installed at Berkeley
 - The system administrator wouldn't permit this
 - He put a deliberate back door into sendmail, to give himself continued access
 - Production systems shipped with this option enabled. . .

Buffer Overflow

- The **finger daemon** call **gets()**, a now-deprecated library routine
- Unlike **fgets()**, there was no buffer length parameter
- By **sending a long-enough string** over the network as input, the attacking program
 1. **Injected some assembler-language code**, and
 2. **Overwrote the return address** in the stack frame so that **gets()** branched to that code instead of back to the caller

Password Guessing

- It looked up a **list of usernames in the password file**
- It used easy transformations of the login name and the user's name, plus a **dictionary of common passwords**

Pre-Authenticated Login

- Exploit trust patterns: /etc/hosts.equiv and per-user .rhosts files list **trusted machines**
- If machine A trusts machine B (if only for a particular user), machine B usually trusts machine A
- This provided two things: **an infection path and a list of other machines to attack**

Spread Patterns

- It looked at a variety of sources to find other machines to attack:
 - rsh/rlogin **trust sources**
 - Machines listed in **.forward files**
 - **Routers** (in 1988, most routers were general-purpose computers)
 - **Randomly-generated addresses on neighboring nets**

Hiding

- The worm used a variety of techniques to hide
- It was named **sh**
- It **forked frequently**, to change processID
- Text **strings** were (lightly) **encrypted**

Essential Elements

- Self-spreading, via **buggy code**
- Self-spreading, via **trust patterns**
- Combination of **directed and random targets** for next attack
- **Stealth** characteristics

Modern Worms

Modern Worms

- Most resemble either the Christmas card worm or the Internet worm
- Today's email worms try to trick the user with tempting Subject: lines
 - million dollar award, software “updates”, etc.
 - A notable one: “Osama bin Laden Captured”, with an attached “video”
 - Some pose as anti-virus software updates. . .
- Can get through many firewalls

Stealthiness

- Deceptive filenames for the attachments
- Add a phony extension before the real one: Saddam_Capture.jpg.exe
- Hide in an encrypted .zip file, with the password in the body of the email
- Many strategies for hiding on hosts, including strange filenames, etc.

Trust Patterns

- Preferentially **attack within the same network** — may be on the inside of a firewall
- Exploit **shared disks**
- Mass-**mailing worms** rely on apparent **trustworthy source**

Spreading Via Buggy Code

- Exploit many different (Windows) bugs
- Can spread much more quickly
- **Slammer** spread about as far as it could in just 15 minutes, and clogged much of the Internet

The Slammer Worm

- Exploited a **bug in Microsoft's SQL server**
- Used UDP, not TCP — a single **376-byte packet to UDP port 1434** could infect a machine!
- **Use of UDP** instead of TCP let it **spread much faster** — one packet, from a forged source address, instead of a three-way handshake, payload transmission, and close() sequence
- **No direct damage, but it clogged network** links very quickly

The Welch Worm

- Attempted to do good
- Used the same Microsoft RPC bug as the Nachi worm
- Removes certain other worm infections
- Installs Microsoft's fix for the hole
- Deletes itself after January 1, 2004

Was it a Good Idea?

- No — unauthorized
- No — not well-tested
- No — generates a lot of network traffic, more than the worm it was trying to cure

Worm Effects

- Seriously clogged networks
- Slammer affected some ATM and air traffic control networks
- CSX Railroad's signaling network was affected

Sobig.F

- Launched in 2003
- Part of a family of worms
- High-quality code
- Primary purpose: **spamming**
- **Turned infected machines into spambots**
- Marked the **turning point in worm design** — now, it's done for **profit** instead of fun

Updating and control

■ Distributed control

- Each worm has a list of other copies
- Ability to create encrypted communication channels to spread info
- Commands cryptographically signed by author.
- Each worm copy, confirms signature, spreads to other copies and then executes the command

■ Programmatic Updates

- Operating systems allow dynamic code loading
- New encrypted attack modules from Worm author

Worm Spread Patterns

How to Own the internet in your spare time [Staniford02]

Spread Patterns?

- The **faster you spread**, the **less likely a defense** could be put up against you
 - More hosts under your control
- **Millions of hosts --> enormous damage**
 - Distributed DOS
 - Access Sensitive Information
 - Create Confusion and Disruption

Code Red I

- Initial version released July 13, 2001.
- Exploited known bug in **Microsoft IIS Web servers**.
- But: **failure to seed random number generator**.
 - All worms attempted to compromise the same sequence of hosts.
- **Linear spread**, didn't get very far

Code Red I v2

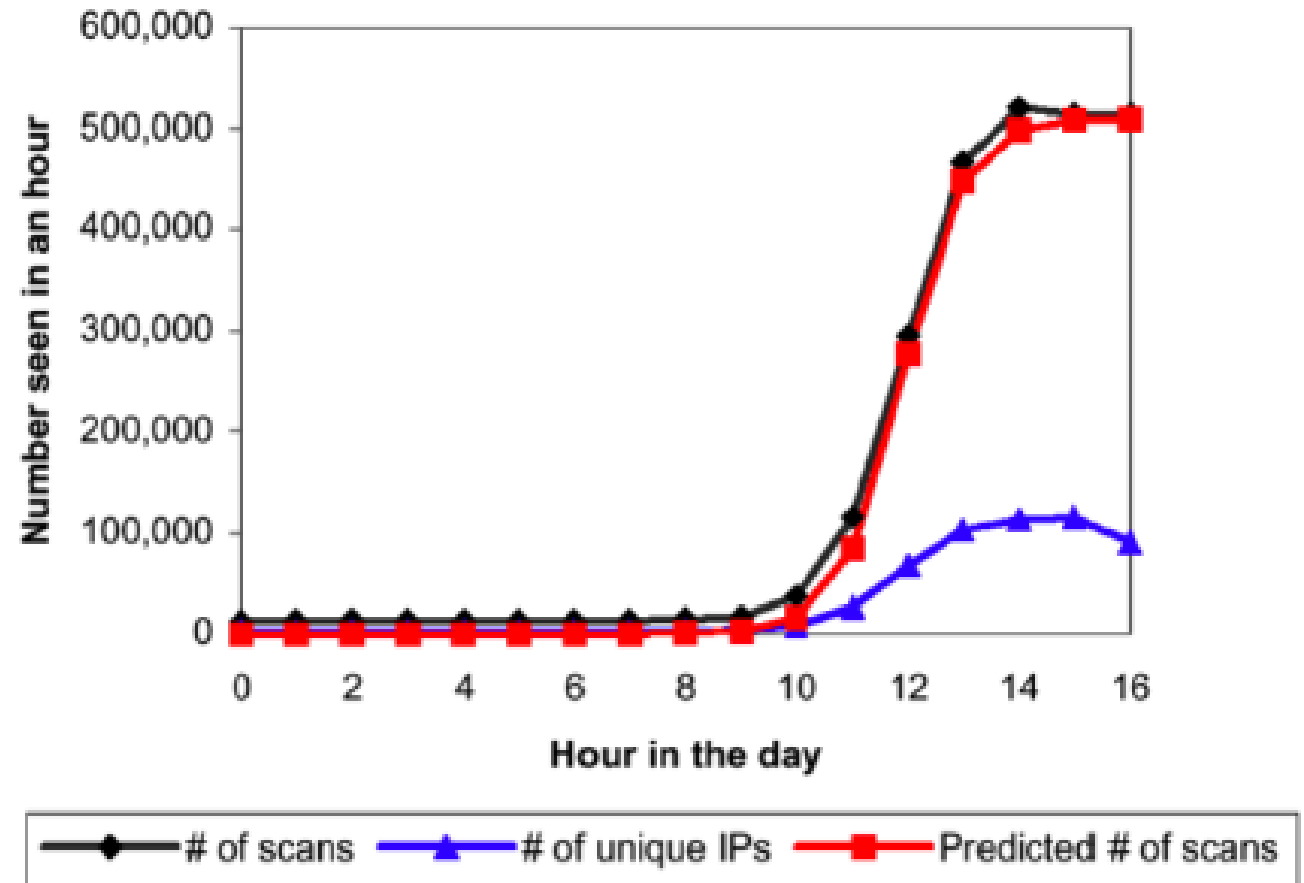
- Released July 19, 2001 (6 days later).
- Same code base but:
 - random number generator correctly seeded.
- DDoS payload targeting IP address of
 - www.whitehouse.gov
- That night, Code Red dies (except for hosts with inaccurate clocks!)

Random Constant Spread Model

- N: Total number of Vulnerable servers in Internet
- K: Initial Compromise Rate: Rate at which a infected host is able to infect new hosts at the start of the incident
- a: Proportion of machines already compromised
- Modeling --> Random Constant Spread (RCS)
 - Gives an exponential equation
 - Depends only on K, not N
- Good enough model (Works for Code Red I)

Hourly

- Hourly probe rate for inbound port 80 at CAS during initial outbreak
- # of unique IPs lags due to the time it takes from getting infected to starting scanning for other victims



Code Red II

- Released August 4, 2001.
- Comment in code: “Code Red II.” But in fact **completely different code base**.
- Payload: a **root backdoor allowing unrestricted remote access**
- Bug: **crashes NT, only works** right on **Windows 2000**.
- Used **localized scanning strategy**

Localized Scanning


- Attempt to infect addresses close to it
 - With probability $\frac{3}{8}$ it chooses a random IP from with the class B (/16) address space of the infected machine
 - With probability $\frac{1}{2}$ from class A (/8)
 - And with probability $\frac{1}{8}$ from the whole internet
- Localized spreading works - hosts around it are often similar, topologically faster, spreads fast in internal network once it gets through the firewall

Nimda

- Released September 18, 2001.
- Multi-mode spreading:
 - attack IIS servers via infected clients.
 - email itself to address book as a virus
 - copy itself across open network shares
 - modifying Web pages on infected servers in order to infect clients
 - scanning for Code Red II and sadmind backdoors (!)

Onset

- Very rapid onset
- Mail based spread → very effective
- Full functionality → ?
- HTTP connections seen at the Lawrence Berkley National Laboratory



Ways of reducing spreading time

- Hit List scanning
- Permutation scanning
- Topological Scanning
- Internet scale hit-lists

Creating Better Worms

■ Hit List Scanning

- “getting off the ground” very fast
 - Say first 10,000 hosts
- Pre-select 10,000-50,000 potentially vulnerable machines
- First worm carries the entire hit list
- Hit list split in half on each infection
- Can establish itself in few seconds

Ways to get Hit list

- Distributed Scanning - use zombies
- Stealthy Scan - spread it over several months
- DNS searches - e. g., www.domain.com
- Spiders - ask the search engines
- Just Listening - P2P,
- or exploit existing worms


Permutation Scanning

- Random scanning inefficient --> **lot of overlap**
 - All worms share a **common pseudo – random permutation**



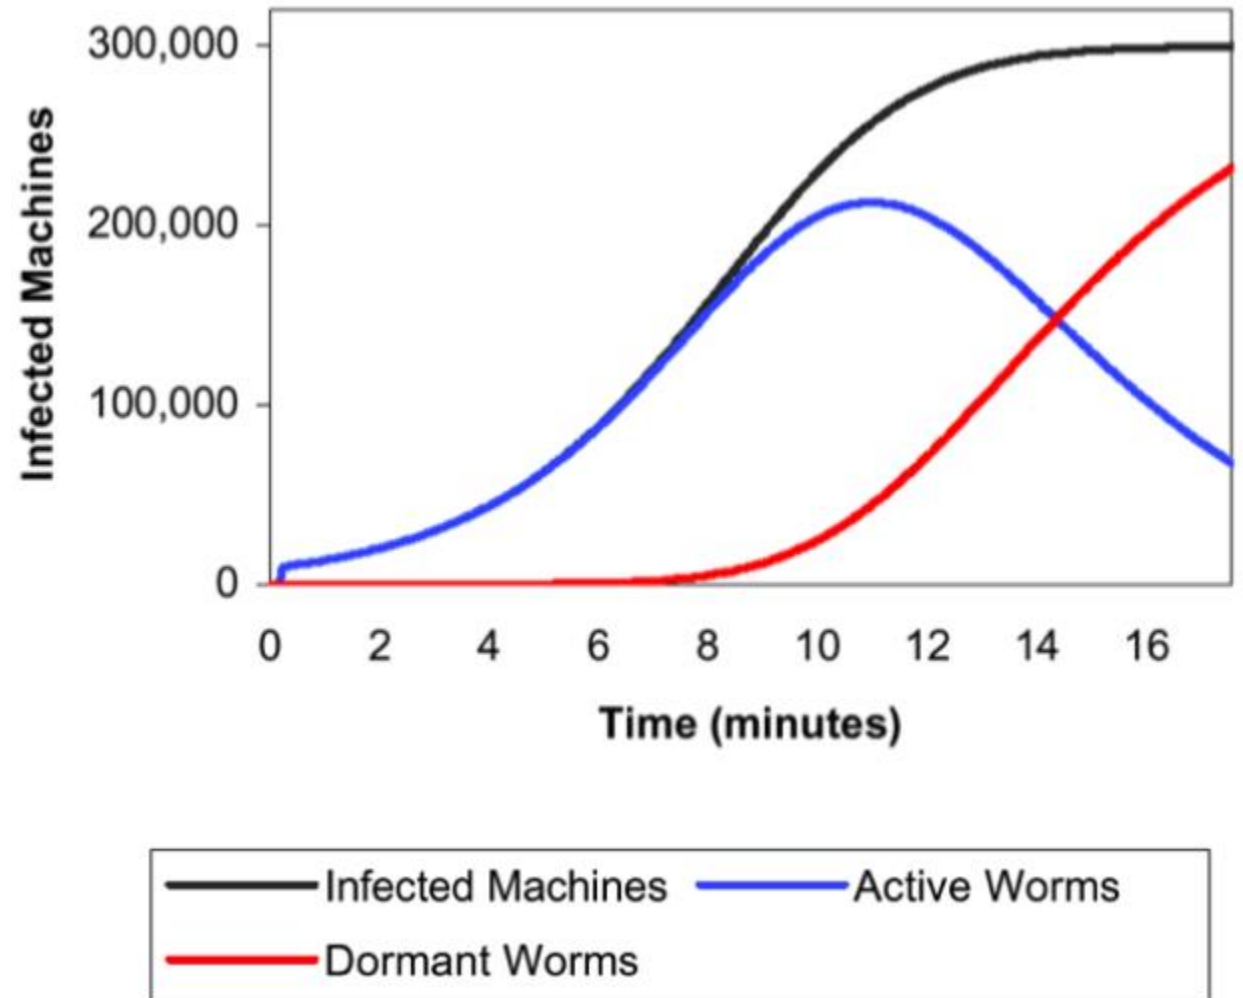
A Warhol Worm

- **Combination of hit-list and permutation scanning**
 - Can spread widely in less than 15 mins
- **Simulation results**
 - 300000 vulnerable machines
 - Conventional (code red) 10 scans/s
 - Fast scanning 100 scans/s
 - Warhol scanning 100 scans/s using 10000 hit-list



Warhol Worm

- Each worm stops when it finds two infected machines without finding any new target
- Rapid growth initially as all worms are active
- As infection nears 100%, many worms go dormant concluding there are few vulnerable machines left



Topological Scanning

- Alternative to hit-list scanning
 - Use addresses available on victim's machines.
 - Email Lists
 - List of web servers from Bookmark
- Use this as a start point before using Permutation Scanning.
- Peer to peer systems are highly vulnerable to this kind of scanning

Faster Worms : Recap


- Fast Startup --> Hit List Scanning
- Extremely Efficient --> Permutation scanning
- Combine the above --> Warhol worms
- Exploit local information --> Topological scanning

Flash Worms

- Fastest Method --> Entire internet in 10s of seconds
- Obtain hit-list of vulnerable servers in advance
 - 2 hours for entire IPv4 space on OC-12 link (622 mbps)
 - List would be big (~ 48 MB)
- Divide into n blocks
 - Infect first of each block and hand over the block to the new worm
 - Repeat for each block
- Alternative: Store pre-assigned chunks on a high BW server
- Two limitations
 - Large list size
 - Latency

Flash Worms

- For 3 million hosts, just 7 layers deep ($n = 10$)




Still need better worms

- All those worms use singular communication patterns
- This forms the basis for automatic detection
- How can we remove that weakness from worms?

Contagion Worms

- Suppose you have two exploits:
 - Es : exploit in web server
 - Ec: exploit in client
- You infect a server (or client) with Es (Ec)
 - When vulnerable client arrives, infect it.
- As client happens to visit other vulnerable servers infects
- Clearly **there are no unusual communication patterns** to be observed (other than slightly larger- than- usual transfers)

Contagion Worms



Contagion Worms

- They become Dangerous with P2P systems because
 - Likely **only need a single exploit**, not a pair.
 - Often, **peers running identical software**.
 - Often used to **transfer large files**.

Contagion Worms

- KaZaA: 9 million distinct IP connections with university hosts (5800) in a single month
- If you Own'd a single university
 - then in November, 2001 you could have Own'd 9 million additional hosts.

References

- [Kapantaidakis] CS588, Giannis Kapantaidakis, University of Crete.
- [Gupta04] Network Security, Ashish Gupta, April 2004.
- [Staniford02]]How to Own the Internet in Your Spare Time, Stuart Staniford, Vern Paxson, and Nicholas Weaver, 11th USENIX Security Symposium (Security 02)
- Steven Bellovin, COMS W4180, Columbia University, 2006
- Mehdi Kharrazi, CE40-817, Sharif University of Technology, 2015
- Vitaly Shmatikov, CS 361S, UT Austin, 2014
- Vitaly Shmatikov, CS 5435, Cornell University, 2022