



# Advanced Network Security

## Worms and Viruses

Amir Mahdi Sadeghzadeh, Ph.D.

# Worms and Viruses

- What are they?
- How do they spread?
- What can be done about them?

# Viruses

- “Infected” program (or floppy)
- When program is executed, it performs its normal function
- It also infects some other programs
- It may carry an extra “payload” that performs other functions

# Worms

- Similar to viruses, but they spread between machines
- Some are fully automatic; some require manual intervention to spread
- Some exploit bugs; others use social engineering

# Classic Worms

# Early Worms

- IBM Christmas Card “Virus”, December 1987
- Morris Internet Worm, November 1988
- Most worms since then have emulated one or both of those

# Modern Worms

- Most resemble either the Christmas card worm or the Internet worm
- Today's email worms try to trick the user with tempting Subject: lines
  - million dollar award, software “updates”, etc.
  - A notable one: “Osama bin Laden Captured”, with an attached “video”
  - Some pose as anti-virus software updates. . .
- Can get through many firewalls

# Stealthiness

- Deceptive filenames for the attachments
- Add a phony extension before the real one: Saddam\_Capture.jpg.exe
- Hide in an encrypted .zip file, with the password in the body of the email
- Many strategies for hiding on hosts, including strange filenames, etc.



# Trust Patterns

- Preferentially **attack within the same network** — may be on the inside of a firewall
- Exploit **shared disks**
- Mass-**mailing worms** rely on apparent **trustworthy source**

# Spreading Via Buggy Code

- Exploit many different (Windows) bugs
- Can spread much more quickly
- **Slammer** spread about as far as it could in just 15 minutes, and clogged much of the Internet

# The Slammer Worm

- Exploited a **bug in Microsoft's SQL server**
- Used UDP, not TCP — a single **376-byte packet to UDP port 1434** could infect a machine!
- **Use of UDP** instead of TCP let it **spread much faster** — one packet, from a forged source address, instead of a three-way handshake, payload transmission, and close() sequence
- **No direct damage, but it clogged network** links very quickly

# Worm Effects

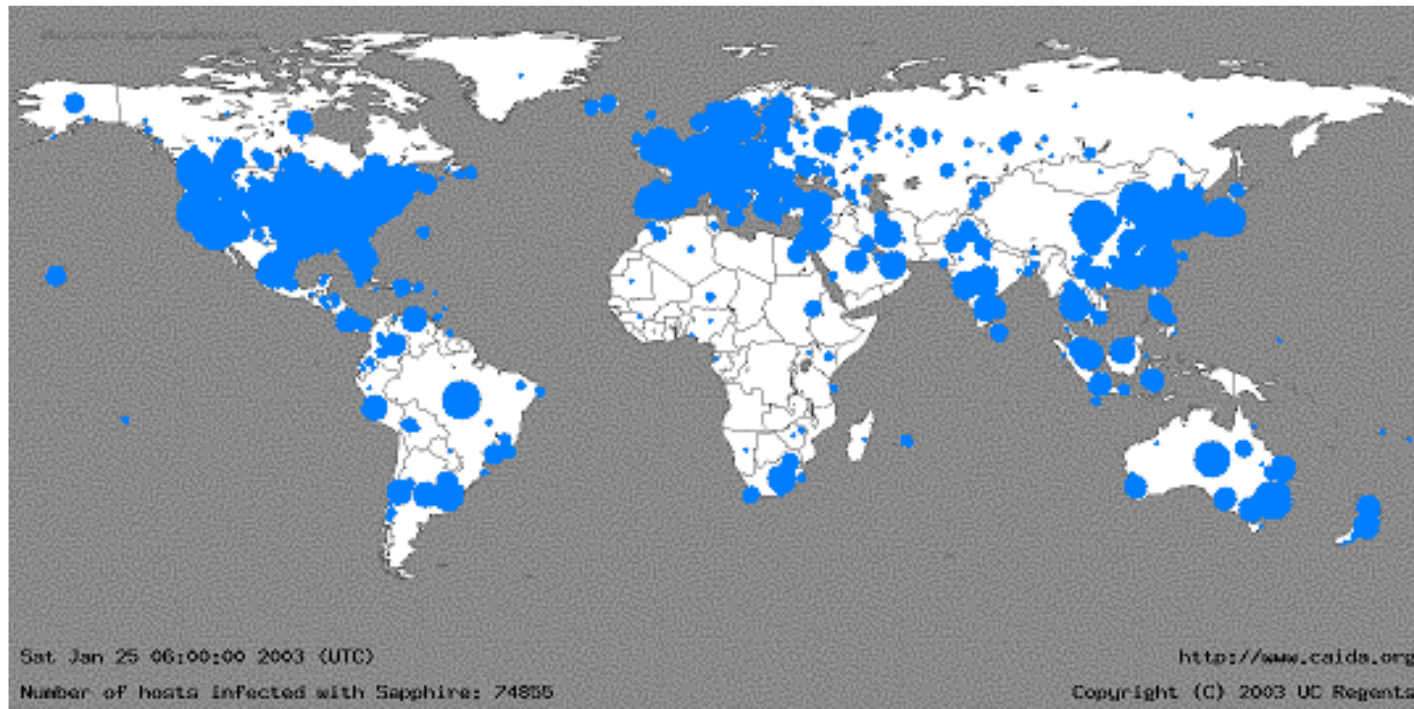
- Seriously clogged networks
- Slammer affected some ATM and air traffic control networks
- CSX Railroad's signaling network was affected

# Problem

- how to **react quickly to worms?**
- CodeRed 2001
  - Infected ~360,000 hosts within 11 hours
- Sapphire/Slammer (376 bytes) 2002
  - Infected ~75,000 hosts within 10 minutes

# The SQL Slammer Worm: 30 Minutes After “Release”

- Infections **doubled every 8.5 seconds**
- Spread 100X faster than Code Red
- At peak, **scanned 55 million hosts per second.**



# Network Effects Of The SQL Slammer Worm

- At the height of infections
  - Several ISPs noted significant bandwidth consumption at peering points
  - Average packet loss approached 20%
  - South Korea lost almost all Internet service for period of time
  - Financial ATMs were affected
  - Some airline ticketing systems overwhelmed

# Current Detection Methods

- Typically an IDS helps the administrators
- Isolation of the worm
- Security experts create the worms signature
- Updates to antivirus and network filtering software
- Correct but expensive, slow and manual procedure.
- Reaction time should be max 60 sec to contain a worm



# Background

- CodeRed in 2001
  - Repair rate: 2% per day - With media attention
  - Automatic Intervention is necessary
- Signature-based models can halt all matching network activity, when the worm's signature is created

# Worm Detection

- Three classes of methods
  - Scan detection
  - Honeypots
  - Behavioral techniques

# Scan Detection

- Look for **unusual frequency and distribution of address scanning**
  - Here is where a **telescope** would be useful
- Limitations
  - Not suited to worms that spread in a **non-random fashion** (i.e. emails, IM, P2P apps)
    - Based on a target list
    - Spread topologically

# Scan Detection

- More limitations
  - Detects infected sites
  - Does not produce a signature

# Honeypots

- Monitored **idle hosts with untreated vulnerabilities**
  - Used to isolate worms
- Limitations
  - **Manual extraction of signatures**
  - **Depend on quick infections**



# Behavioral Detection

- Looks for unusual system call patterns
  - Sending a packet from the same buffer containing a received packet
  - Can detect slow moving worms
- Limitations
  - Needs application-specific knowledge
  - Cannot infer a large-scale outbreak

# Characterization

- Process of analyzing and identifying a new worm
- Current approaches
  - Use a priori vulnerability signatures
  - Automated signature extraction

# Next Step: Vulnerability Signatures

- Example
  - **Slammer Worm**
    - UDP traffic on port 1434 that is longer than 100 bytes (buffer overflow)
- Can be deployed **before the outbreak**
  - **Can only be applied to well-known vulnerabilities**



# Some Automated Signature Extraction Techniques

- Allows worm to infect decoy programs
  - Extracts the modified regions of the decoy
  - Uses heuristics to identify invariant code strings across infected instances

# Some Automated Signature Extraction Techniques

- Limitation

- Assumes the presence of a worm in a controlled environment

# Next Step: Containment

- Mechanism used to deter the spread of an active worm
  - Host quarantine
    - Via IP ACLs on routers or firewalls
- String-matching
- Connection throttling
  - On all outgoing connections

Automated Worm Fingerprinting, Sumeet Singh, Cristian Estan, George Varghese and Stefan Savage, Proceedings of the ACM/USENIX Symposium on Operating System Design and Implementation, San Francisco, CA, December 2004.

# Earlybird

- Automatic detection and containment of new worms
- Content Sifting:
  - Content of worm traffic is invariant
  - Worm spread dynamics atypical of Internet Applications
- Frequently repeated and widely dispersed content strings -> new worm

# Defining Worm Behavior

- **Content invariance**
  - Portions of a worm are invariant (e.g. the decryption routine)
- **Content prevalence**
  - Appears frequently on the network
- **Address dispersion**
  - Distribution of destination addresses more uniform to spread fast

# Finding Worm Signatures

- Traffic pattern is sufficient for detecting worms
  - Relatively straightforward
  - Extract all possible substrings
  - Raise an alarm when
    - `FrequencyCounter[substring] > threshold1`
    - `SourceCounter[substring] > threshold2`
    - `DestCounter[substring] > threshold3`

# Practical Content Sifting

- Characteristics
  - Small processing requirements
  - Small memory requirements
  - Allows arbitrary deployment strategies



# Estimating Content Prevalence


- Finding the packet payloads that appear at least  $x$  times among the  $N$  packets sent
  - During a given interval

# Estimating Content Prevalence


- Given a 1Gbps
  - Table[payload]
    - 1 GB table filled in **less than 10 seconds**
  - Table[hash[payload]]
    - 1 GB table filled in **4 minutes**
    - **Tracking millions of ants to track a few elephants**
    - **Collisions**...false positives

# Multistage Filters


Array of  
counters



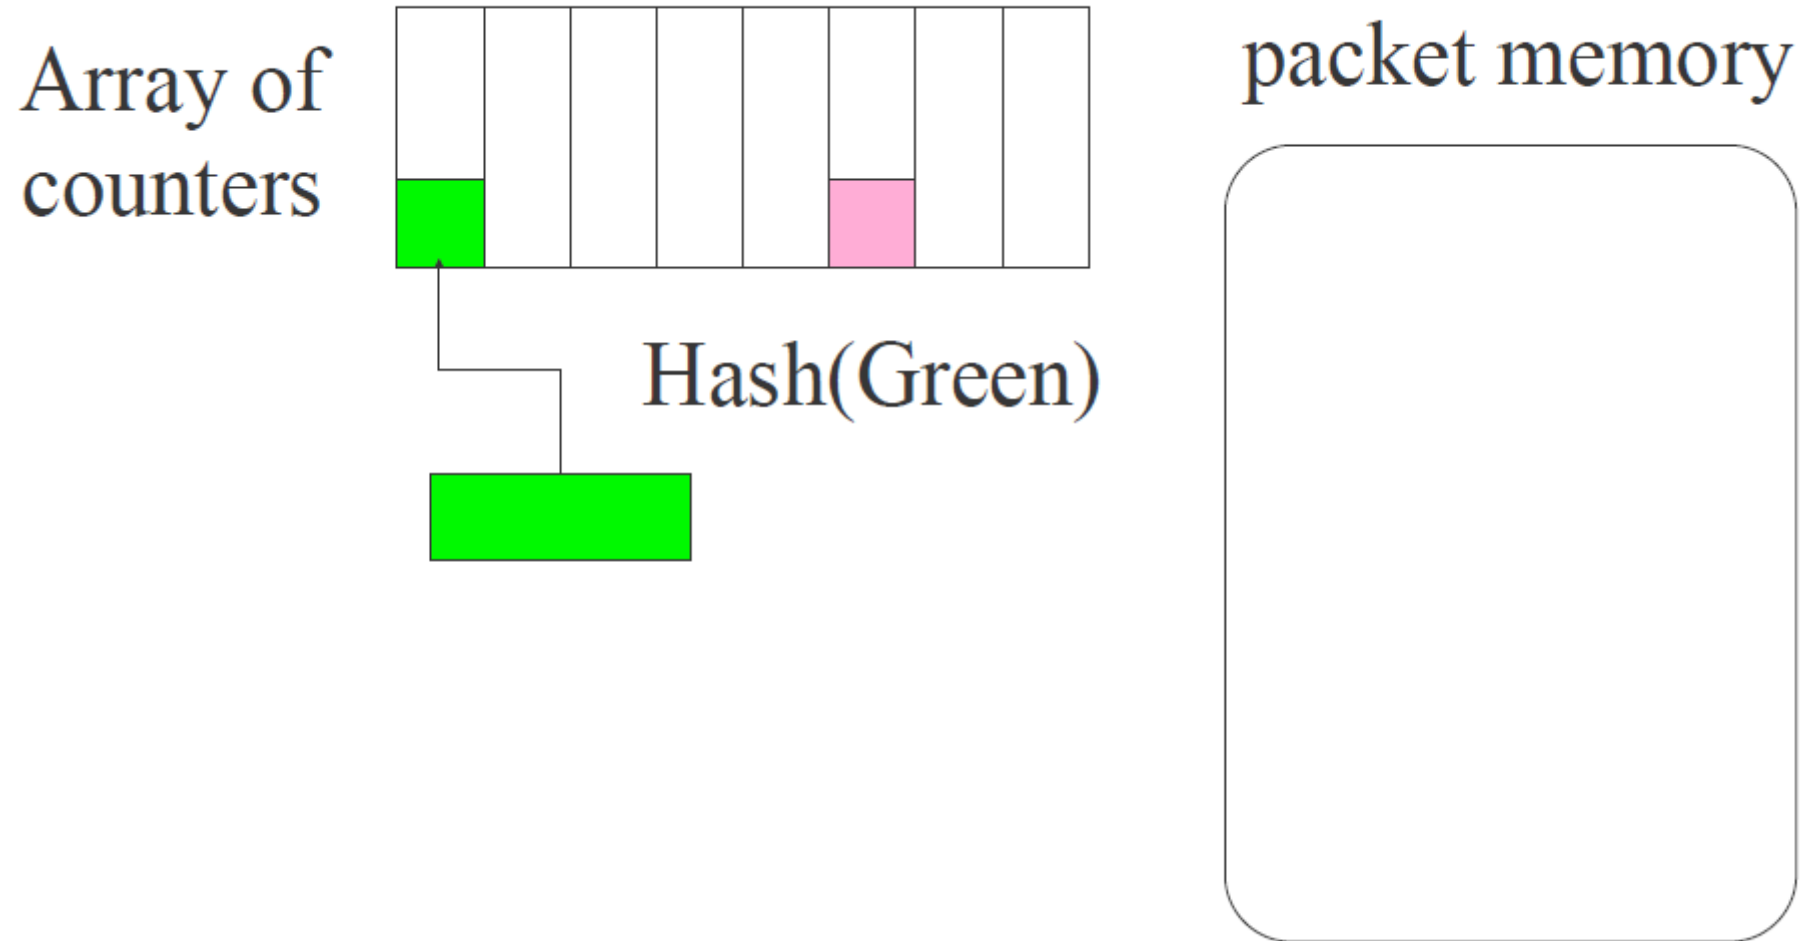
Hash(Pink)




stream memory




# Multistage Filters




# Multistage Filters



# Multistage Filters




packet memory




# Multistage Filters

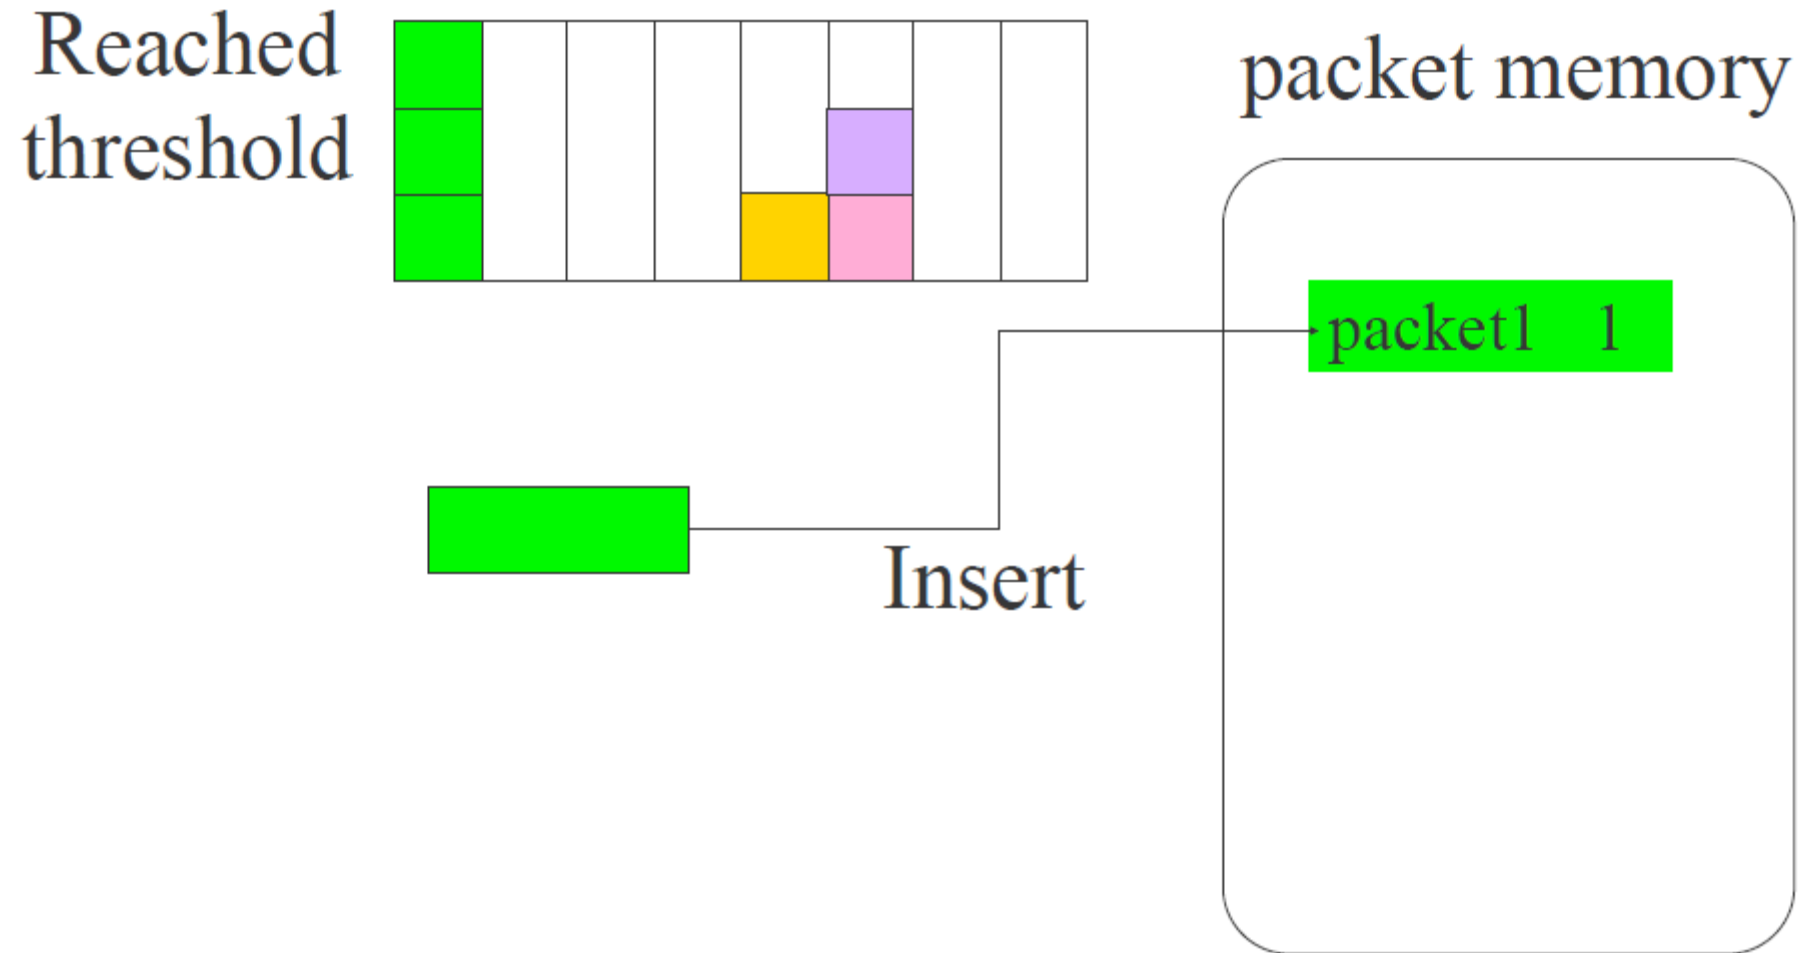
Collisions  
are OK



packet memory




# Multistage Filters





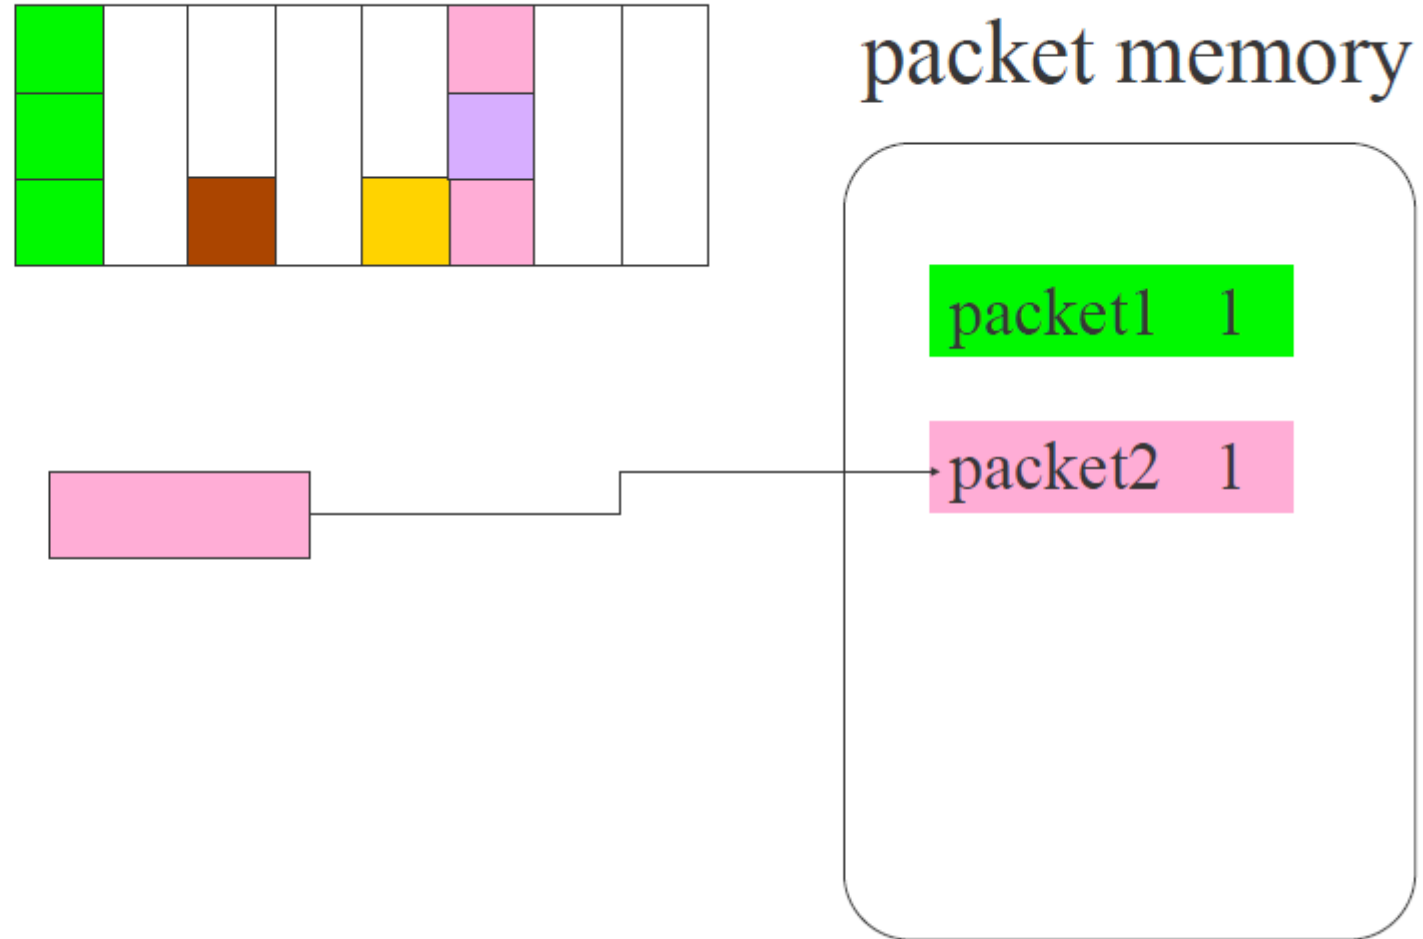
# Multistage Filters



packet memory


packet1 1

# Multistage Filters




# Multistage Filters

Stage 1



Stage 2



packet memory

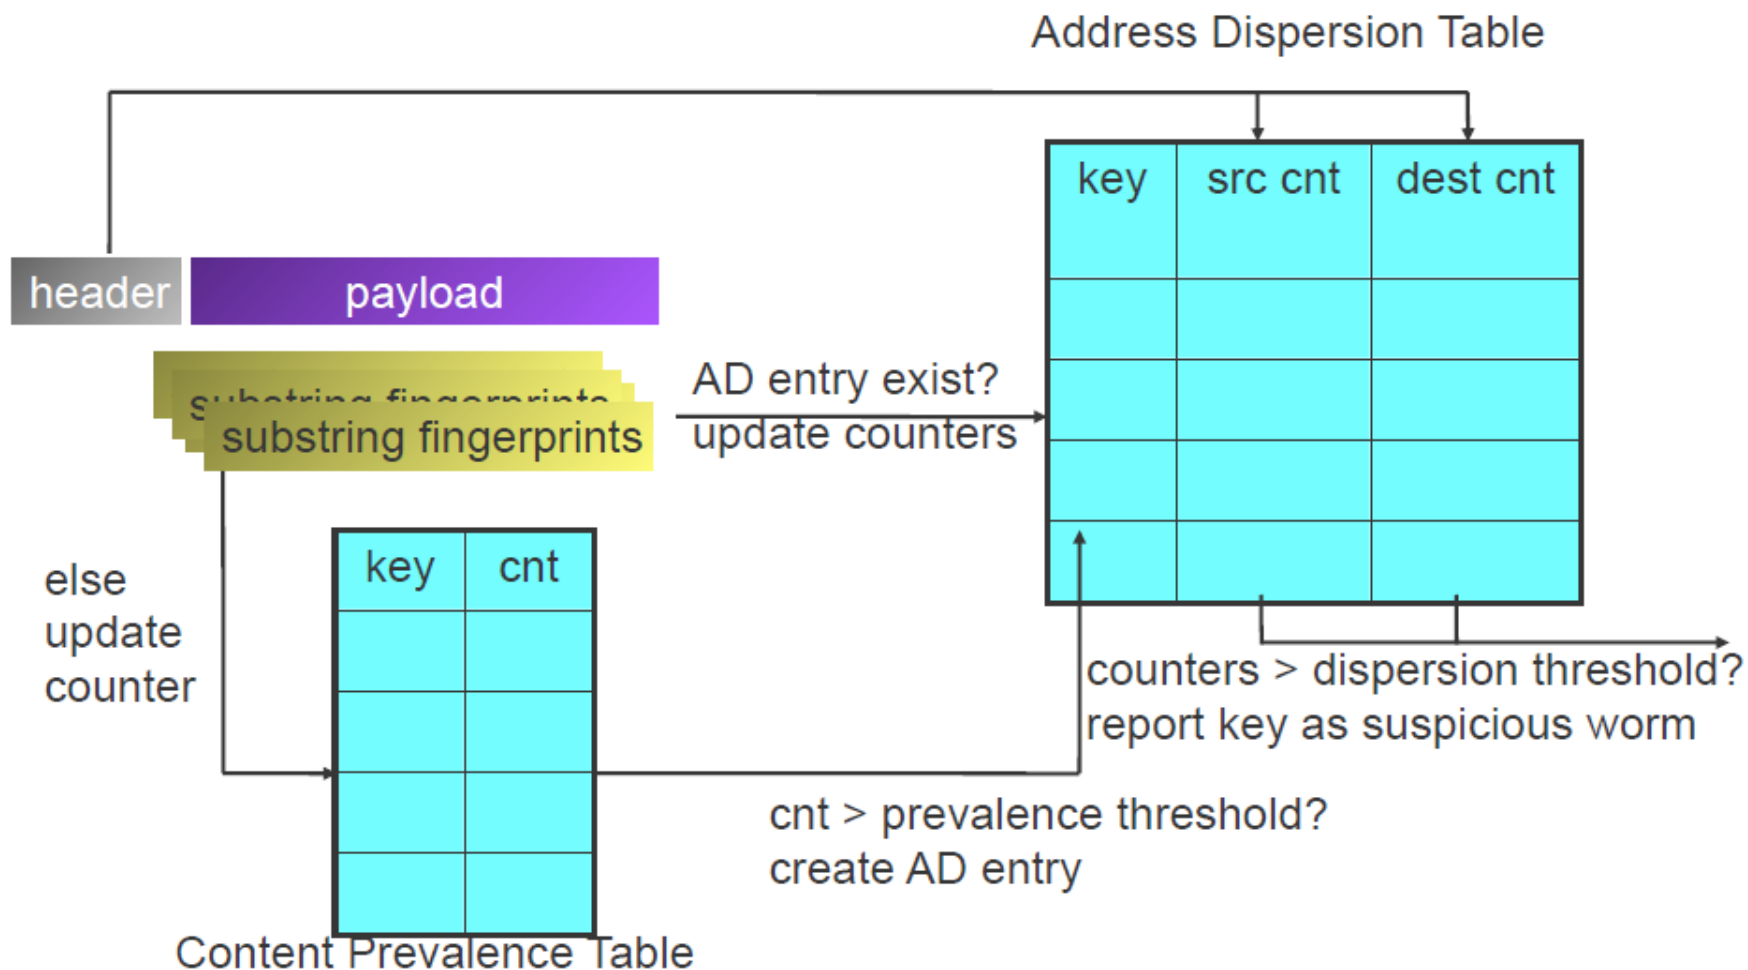
packet1 1

No false negatives!

# Estimating Address Dispersion

- **Not sufficient** to count the number of **source and destination pairs**
  - e.g. send a mail to a mailing list
    - Two sources—mail server and the sender
    - Many destinations
- Need to track the **distinct source and destination IP** addresses
  - **For each substring**
- Simple list or hash table is too expensive
  - Use **Bitmap data structure**

# Putting It Together



# System Design

- Two major components
  - Sensors
    - Sift through traffic for a given address space
    - Report signatures
  - An aggregator
    - Coordinates real-time updates
    - Distributes signatures

# Implementation and Environment

## ■ Implementation

- Written in C and MySQL (5,000 lines)
- rrd-tools library for graphical reporting
- PHP scripting for administrative control
- Prototype executes on a 1.6Ghz AMD Opteron 242 1U Server
  - Linux 2.6 kernel

- Processes **1TB of traffic per day**
- Can **keep up with 200Mbps** of continuous traffic



# Content prevalence threshold

- Using a 60 second measurement interval and a whole packet CRC
  - over 97 percent of all signatures repeat two or fewer times and 94.5 percent are only observed once
- Using a finer grained content hash or a longer measurement interval increases these numbers even further
- Default: 3 repetitions




Figure 8: Cumulative distribution function of content signatures for different hash functions. This CDF is computed from the set of repetitions found in each measurement interval over a period of 10 minutes. Note that the y axis is artificially truncated to show more detail.



# Address dispersion threshold

- After 10 minutes there are over 1000 signatures with a low dispersion threshold of 2
- Using a threshold of 30, there are only 5 or 6 prevalent strings meeting the dispersion criteria
- Default: 30 sources and 30 destinations




Figure 9: Number of distinct signatures detected over time for different address dispersion thresholds.

# Garbage Collection

- When the timeout is set to 100 seconds, then almost 60 percent of all signatures are garbage collected before a subsequent update
- Using a timeout of 1000 seconds, this number is reduced to roughly 20 percent of signatures
- Default: **several hours**




Figure 10: Log-scale cumulative distribution function of the maximum time period between updates for entries in the address dispersion table.

# Performance Processing Time

	Mean	Std. Dev.
<b>Component wise breakdown</b>		
<b>Rabin Fingerprint</b>		
First Fingerprint (40 bytes)	0.349	0.472
Increment (each byte)	0.037	0.004
<b>Multi Stage Filter</b>		
Test & Increment	0.146	0.049
<b>AD Table Entry</b>		
Lookup	0.021	0.032
Update	0.027	0.013
Create	0.252	0.306
Insert	0.113	0.075
<b>Overall Packet</b>		
Header Parsing & First Fingerprint	0.444	0.522
Per-byte processing	0.409	0.148
<b>Overall Packet with Flow-Reassembly</b>		
Header Parsing & Flow maintenance	0.671	0.923
Per-byte processing	0.451	0.186

Table 1: This table shows overhead (in microseconds) incurred by each of the individual operations performed on a packet. The mean and standard deviation are computed over a 10 minute interval (25 million Packets). This table represents raw overheads before sampling. Using 1 in 64 value sampling, the effective mean per byte processing time reduces to 0.042 microseconds.

# Performance Memory Consumption

- Prevalence table
  - Totals to 2 MB
- Address Dispersion Table
  - under 1MB
- Total less than 4MB

# Trace-Based Verification

- Two main **sources of false positives**
  - 2,000 common protocol headers
    - e.g. HTTP, SMTP
    - Whitelisted
  - SPAM e-mails
  - BitTorrent
    - Many-to-many download

# False Negatives

- So far none
- Detected every worm outbreak

# Evasions


- An attacker might evade detection by **splitting an invariant string across packets**
  - Have fingerprints across packets
- **Traffic normalization**
  - remember **attacks on IDS**
- **Polymorphic** viruses
  - Semantically equivalent but **textually distinct code**
  - **Invariant decoding routine**

POLYGRAPH: Automatically Generating Signatures for Polymorphic Worms, James Newsome, Brad Karp, Dawn Song, IEEE Security and Privacy Symposium, May 2005.



# Challenge: Polymorphic Worms

- Polymorphic worms **minimize invariant content**
  - Encrypted payload
  - Obfuscated decryption routine
- Polymorphic tools are already available
  - Clet, ADMmutate



# Challenge: Polymorphic Worms

- Polymorphic worms  
minimize invariant content


- 

- 

Do good signatures for polymorphic worms exist?

- Polymorphic worms  
Can we generate them *automatically*?

- 



# Goals

- Identify **classes of signatures** that can
  - Accurately describe polymorphic worms
  - Be used to filter a high speed network line
  - Be generated automatically and efficiently

# Good News: Still some invariant content

## ■ Protocol framing

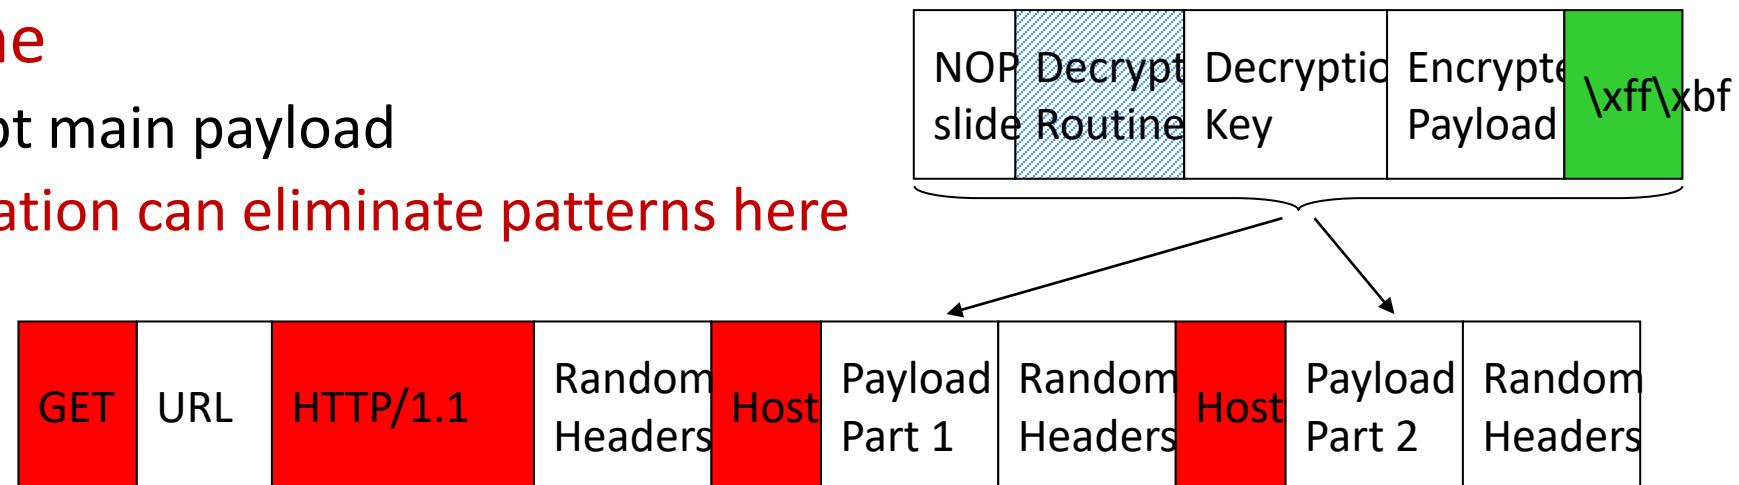
- Needed to make server go down vulnerable code path

## ■ Overwritten Return Address

- Needed to redirect execution to worm code

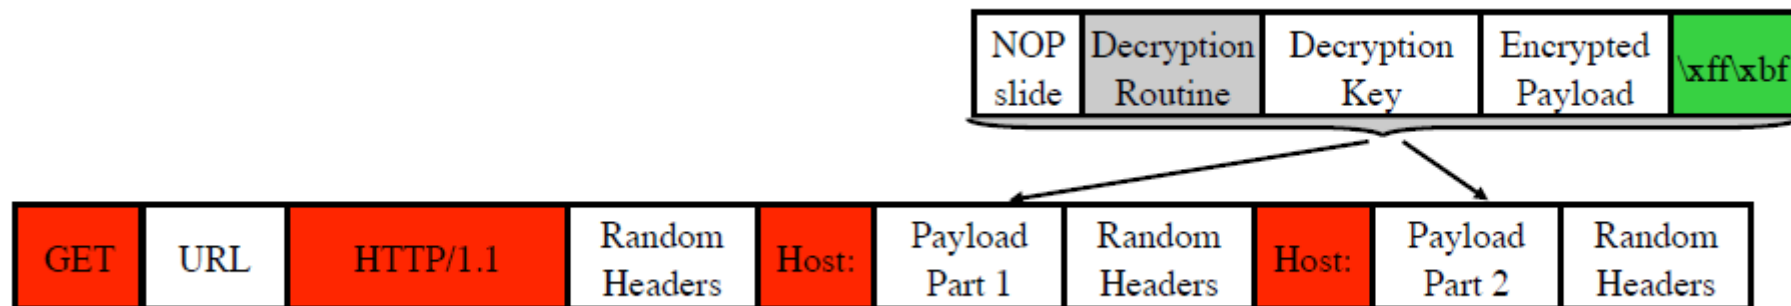
## ■ Decryption routine

- Needed to decrypt main payload
- BUT, code **obfuscation can eliminate patterns here**



# Bad News: Previous Approaches Insufficient

- Previous approaches use a common substring
- Longest substring
  - “HTTP/1.1”
  - 93% false positive rate
- Most specific substring
  - “\xff\xbf”
  - .008% false positive rate (10 / 125,301)

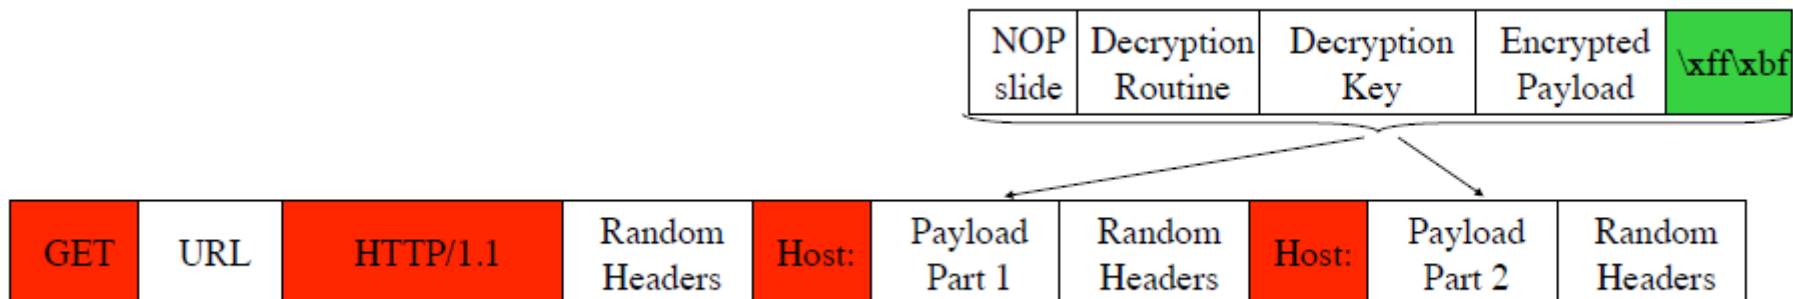


# What to do?

- No one substring is specific enough
- BUT, there are multiple substrings
  - Protocol framing
  - Value used to overwrite return address
  - (Parts of poorly obfuscated code)
- Our approach: combine the substrings

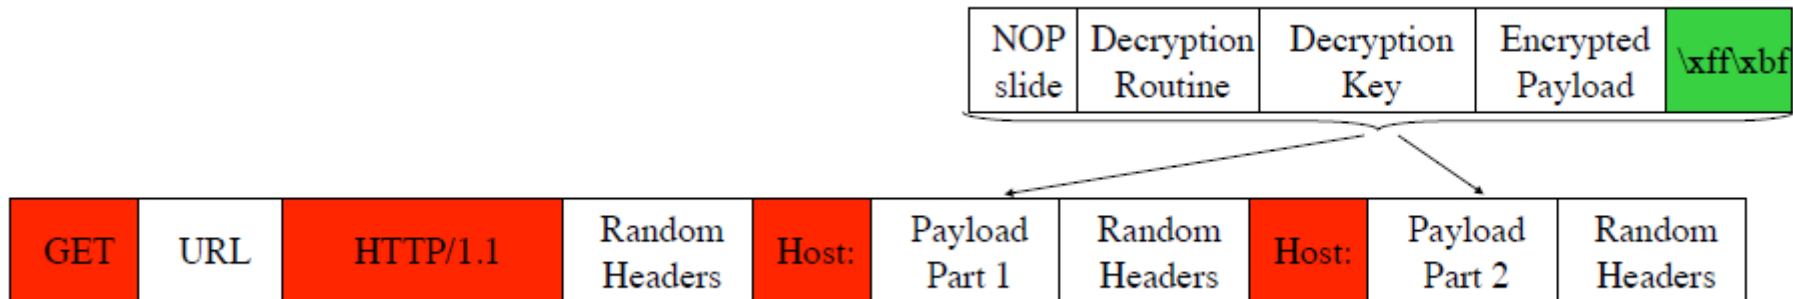
# Signature Class (I): Conjunction

- Signature is a **set of strings** (tokens)
- Flow matches signature iff it **contains all tokens in the signature**
- $O(n)$  time to match ( $n$  is flow length)
- Generated signature:
  - “GET” and “HTTP/1.1” and “\r\nHost:” and “\r\nHost:” and “\xff\xbf”
  - .0024% false positive rate (3 / 125,301)



# Signature Class (II): Token Subsequence

- Signature is an **ordered set of tokens**
- Flow matches iff it contains **all the tokens in signature, in the given order**
- $O(n)$  time to match ( $n$  is flow length)
- Generated signature:
  - GET.\*HTTP/1.1.\*\r\nHost:.\*\r\nHost:.\*\xff\xbf
  - .0008% false positive rate (1 / 125,301)





# Experiment: Signature Generation

- How many worm samples do we need?
  - Too few samples --> signature is too specific --> false negatives
- Experimental setup
  - Using a 15 day port 80 trace from lab perimeter
  - Innocuous pool: First 5 days (45,111 streams)
  - Suspicious Pool:
    - Using Apache exploit described in paper
    - Non-invariant portions filled with random bytes
- Signature evaluation:
  - False positives: Last 10 days (125,301 streams)
  - False negatives: 1000 generated worm samples

# Signature Generation Results

# Worm Samples	Conjunction	Subseq
2	100% FN	100% FN
3 to 100	0% FN .0024% FP	0% FN .0008% FP

GET .\* HTTP/1.1\r\n.\*\r\nHost: .\*\xee\x7b.\*\xb2\xe.\*\r\nHost: .\*\xef  
.\*\xa3.\*\x8b\xf4.\*\x89\x8b.\*E\xeb.\*\xff\xbf

GET .\* HTTP/1.1\r\n.\*\r\nHost: .\*\r\nHost:.\*\xff\xbf

# Acknowledgments/References

- [Singh04] Automated Worm Fingerprinting, Sumeet Singh, Cristian Estan, George Varghese and Stefan Savage, Proceedings of the ACM/USENIX Symposium on Operating System Design and Implementation, San Francisco, CA, December 2004.
- [Wang05] FSU COP 5611 (Spring 2005) Advanced Operating Systems by Andy Wang.
- [kuzma] [ww.cs.northwestern.edu/~akuzma/classes/CS495-s05/doc/awf.ppt](http://ww.cs.northwestern.edu/~akuzma/classes/CS495-s05/doc/awf.ppt)
- [hy558] [ww.csd.uoc.gr/~hy558/reports/itsomp\\_fingerprinting.ppt](http://ww.csd.uoc.gr/~hy558/reports/itsomp_fingerprinting.ppt)
- [zou07] Research in Computer and Network Security, CDA6938, Cliff Zou, 2007
- [Newsome05] POLYGRAPH: Automatically Generating Signatures for Polymorphic Worms, James Newsome, Brad Karp, Dawn Song, presentation at IEEE Security and Privacy Symposium, May 2005.